



1. What is a Process?

- **⋄** A process is an executing program.
- **A process:**

\$ cat file1 file2 &

Two processes:

\$ ls | wc - 1

*Each user can run many processes
at once
(e.g. using &)



A More Precise Definition

- A process is the context (the information/ data) maintained for an executing program.
- Intuitively, a process is the abstraction of a physical processor.
 - Exists because it is difficult for the OS to otherwise coordinate many concurrent activities, such as incoming network data, multiple users, etc.
- IMPORTANT: A process is sequential





What makes up a Process?

- program code
- machine registers
- stack
- open files (file descriptors)
- *an environment (environment variables; credentials for security)





Some of the Context Information

Process ID (pid)

Parent process ID (ppid)

Real User ID which

Effective User ID

- Current directory
- File descriptor table
- Environment

unique integer

ID of user/process started this process

ID of user who wrote the process' program

VAR=VALUE pairs

continued



Pointer to program code

Pointer to dataMemory for global vars

Pointer to stackMemory for local vars

Pointer to heap Dynamically allocated

Execution priority

Signal information

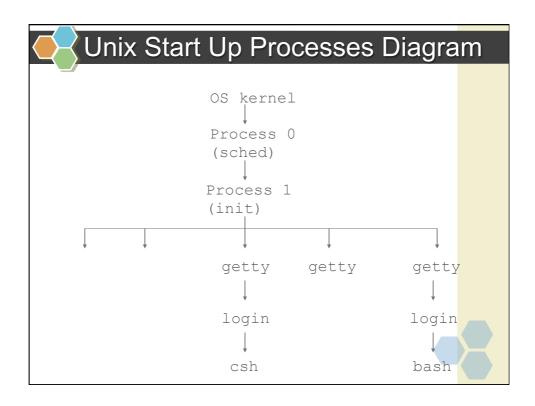




Important System Processes

- *init Mother of all processes. init is started at boot time and is responsible for starting other processes.
 - init uses file inittab & directories: /etc/rc?.d
- *getty login process that manages login sessions.







Pid and Parentage

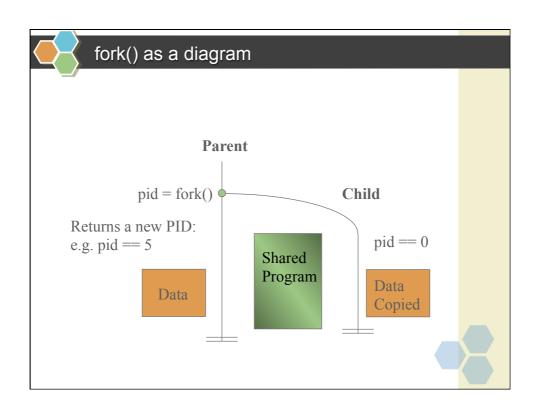
- ❖ A process ID or *pid* is a positive integer that uniquely identifies a running process, and is stored in a variable of type *pid_t*.
- You can get the process pid or parent's pid

```
#include <sys/types>
main()
{
  pid_t pid, ppid;
  printf( "My PID is:%d\n\n",(pid = getpid()) );
  printf( "Par PID is:%d\n\n",(ppid = getppid()) );
}
```



2. fork()

- #include <sys/types.h>
 #include <unistd.h>
 pid t fork(void);
- Creates a child process by making a copy of the parent process --- an exact duplicate.
 - Implicitly specifies code, registers, stack, data, files
- ❖ Both the child and the parent continue running.





Process IDs (pids revisited)

- pid = fork();
- In the child: pid == 0;
 In the parent: pid == the process ID of
 the child.
- *A program almost always uses this pid difference to do different things in the parent and child.



fork() Example (parchld.c)



```
CHILD 0
CHILD 1
CHILD 2

PARENT 0
PARENT 1
PARENT 2
PARENT 3
CHILD 3
CHILD 4

PARENT 4
:
```



Things to Note

- i is copied between parent and child.
- The switching between the parent and child depends on many factors:
 - machine load, system process scheduling
- ❖I/O buffering effects amount of output shown.
- Output interleaving is nondeterministic
 - cannot determine output by looking at code





3. exec()

*Family of functions for replacing process's program with the one inside the exec() call. e.g.

Same as "sort -n foobar"

tinymenu.c

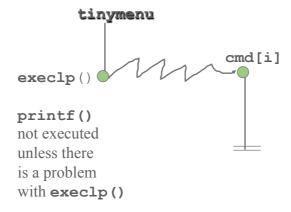
```
#include <stdio.h>
#include <unistd.h>

void main()
{ char *cmd[] = {"who", "ls", "date"};
  int i;
  printf("0=who 1=ls 2=date : ");
  scanf("%d", &i);

  execlp( cmd[i], cmd[i], (char *)0 );
  printf( "execlp failed\n" );
}
```



Execution





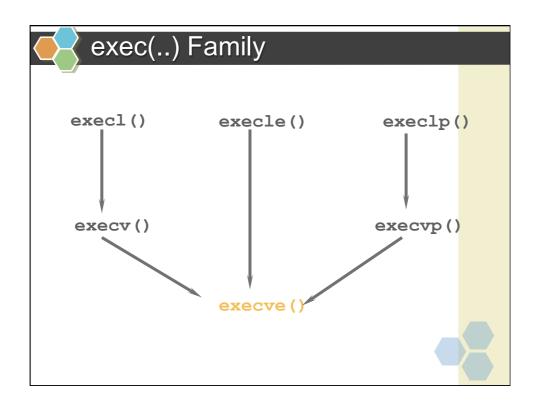


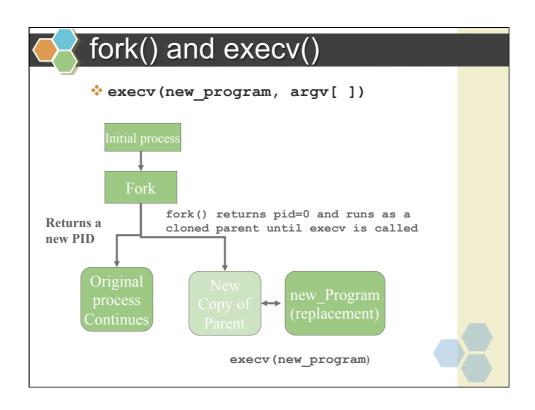
exec(..) Family

There are 6 versions of the exec function, and they all do about the same thing: they replace the current program with the text of the new program. Main difference is how parameters are passed.











4. wait()

- #include <sys/types.h>
 #include <sys/wait.h>
 pid t wait(int *statloc);
- Suspends calling process until child has finished. Returns the process ID of the terminated child if ok, -1 on error.
- * statloc can be (int *)0 or a variable which will be bound to status info. about the child.



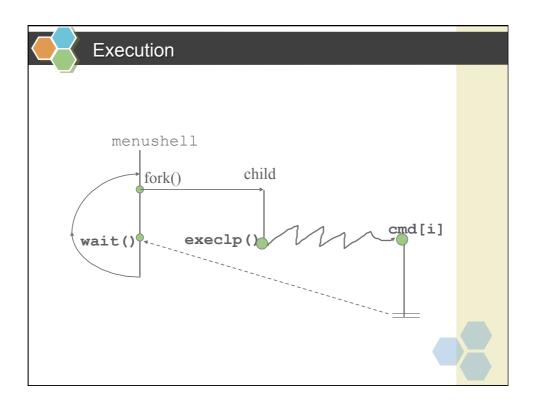
wait() Actions

- ❖A process that calls wait() can:
 - suspend (block) if all of its children are still running, or
 - return immediately with the termination status of a child, or
 - return immediately with an error if there are no child processes.



```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

void main()
{
    char *cmd[] = {"who", "ls", "date"};
    int i;
    while(1)
        {
        printf(0=who 1=ls 2=date : ");
        scanf("%d", &i);
        :
        continued
```





Macros for wait (1)

- **❖** WIFEXITED(status)
 - Returns true if the child exited normally.
- **❖** WEXITSTATUS(*status*)
 - Evaluates to the least significant eight bits of the return code of the child which terminated, which may have been set as the argument to a call to exit() or as the argument for a return.
 - This macro can only be evaluated if WIFEXITED returned non-zero.





Macros for wait (2)

❖ WIFSIGNALED(*status*)

 Returns true if the child process exited because of a signal which was not caught.

WTERMSIG(status)

- Returns the signal number that caused the child process to terminate.
- This macro can only be evaluated if WIFSIGNALED returned non-zero.



waitpid()

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t waitpid( pid_t pid, int *status, int opts )
```

❖ waitpid - can wait for a particular child

❖ pid < -1</p>

 Wait for any child process whose process group ID is equal to the absolute value of pid.

pid == -1

- Wait for any child process.
- Same behavior which wait() exhibits.
- pid == 0
- Wait for any child process whose process group ID is equal to that of the calling process.



🍫 pid > 0

- Wait for the child whose process ID is equal to the value of pid.
- options
 - Zero or more of the following constants can be ORed.
 - WNOHANG
 - » Return immediately if no child has exited.
 - WUNTRACED
 - » Also return for children which are stopped, and whose status has not been reported (because of signal).
- Return value
 - The process ID of the child which exited.
 - -1 on error; 0 if WNOHANG was used and no child was available.



Macros for waitpid

❖ WIFSTOPPED(status)

- Returns true if the child process which caused the return is currently stopped.
- This is only possible if the call was done using WUNTRACED.

WSTOPSIG(status)

- Returns the signal number which caused the child to stop.
- This macro can only be evaluated if WIFSTOPPED returned non-zero.



```
#include <stdio.h>
#include <sys/wait.h>
#include <sys/types.h>

int main(void)
{
    pid_t pid;
    int status;

if( (pid = fork() ) == 0 )
    { /* child */
    printf("I am a child with pid = %d\n",
        getpid());
        sleep(60);
        printf("child terminates\n");
        exit(0);
    }
```

```
{ /* parent */
       while (1)
           waitpid( pid, &status, WUNTRACED );
           if( WIFSTOPPED(status) )
           {
                  printf("child stopped, signal(%d)\n'',
                          WSTOPSIG(status));
                  continue;
           else if( WIFEXITED(status) )
                  printf("normal termination with
                      status(%d)\n",
                       WEXITSTATUS(status));
           else if (WIFSIGNALED(status))
                  printf("abnormal termination,
                       signal(%d)\n'',
                       WTERMSIG(status));
           exit(0);
           } /* while */
   } /* parent */
} /* main */
```



5. Process Data

- Since a child process is a copy of the parent, it has copies of the parent's data.
- *A change to a variable in the child will *not* change that variable in the parent.



Example

(globex.c)

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int globvar = 6;
char buf[] = "stdout write\n";

int main(void)
{
   int w = 88;
   pid_t pid;
```

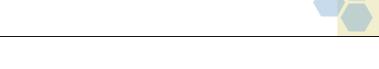


```
Output
  $ globex
                  /* write not buffered */
  stdout write
 Before fork()
  pid = 430, globvar = 7, w = 89
                    /*child chg*/
  pid = 429, globvar = 6, w = 88
                    /* parent no chg */
$ $ globex > temp.out
  $ cat temp.out
  stdout write
  Before fork()
  pid = 430, globvar = 7, w = 89
  Before fork() /* fully buffered */
  pid = 429, qlobvar = 6, w = 88
```



6. Process File Descriptors

- A child and parent have copies of the file descriptors, but the R-W pointer is maintained by the system:
 - the R-W pointer is shared
- This means that a read() or write() in one process will affect the other process since the R-W pointer is changed.



Example: File used across processes

(shfile.c)

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <unistd.h>
#include <fcntl.h>
void printpos(char *msg, int fd);
void fatal(char *msg);

int main(void)
{ int fd;    /* file descriptor */
    pid_t pid;
    char buf[10]; /* for file data */
```



```
if ((fd=open("data-file", O_RDONLY)) < 0)
    perror("open");

    read(fd, buf, 10); /* move R-W ptr */

    printpos( "Before fork", fd );

    if( (pid = fork()) == 0 )
        { /* child */
        printpos( "Child before read", fd );
        read( fd, buf, 10 );
        printpos( " Child after read", fd );
    }
    :</pre>
```

```
void printpos( char *msg, int fd )
/* Print position in file */
    {
    long int pos;

    if( (pos = lseek( fd, OL, SEEK_CUR) ) < OL )
        perror("lseek");

    printf( "%s: %ld\n", msg, pos );
}</pre>
```

```
$ shfile

Before fork: 10
Child before read: 10
Child after read: 20
Parent after wait: 20

what's happened?
```



8. Special Exit Cases

Two special cases:

- *1) A child exits when its parent is not currently executing wait()
 - the child becomes a zombie
 - status data about the child is stored until the parent does a wait()





- 2) A parent exits when 1 or more children are still running
 - children are adopted by the system's initialization process (/etc/init)
 - it can then monitor/kill them





9. I/O redirection

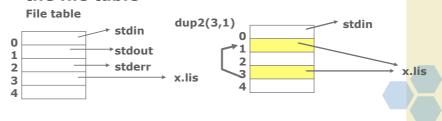
The trick: you can change where the standard I/O streams are going/ coming from after the fork but before the exec



Re

Redirection of standard output

- **⋄** Example implement shell: Is > x.ls
- program:
 - Open a new file x.lis
 - Redirect standard output to x.lis using dup command
 everything sent to standard output ends in x.lis
 - execute Is in the process
- dup2(int fin, int fout) copies fin to fout in the file table



Example - implement ls > x.lis

```
#include <unistd.h>
int main ()
    {
    int fileId;
    fileId = creat( "x.lis",0640 );
    if( fileId < 0 )
        {
        printf( stderr, "error creating x.lis\n" );
        exit (1);
        }
    dup2( fileId, stdout ); /* copy fileID to stdout */
    close( fileId );
    execl( "/bin/ls", "ls", 0 );
    }
}</pre>
```



10. User and Group ID

♦ Group ID

Real. effective

User ID

- Real user ID
 - Identifies the user who is responsible for the running process.
- Effective user ID
 - Used to assign ownership of newly created files, to check file access permissions, and to check permission to send signals to processes.
 - To change euid: executes a setuid-program that has the setuid bit set or invokes the setuid() system call.
 - The setuid(uid) system call: if euid is not superuser, uid must be the real uid or the saved uid (the kernel also resets euid to uid).
- Real and effective uid: inherit (fork), maintain (exec).



Read IDs

- pid_t getuid(void);
 - Returns the real user ID of the current process
- pid_t geteuid(void);
 - Returns the effective user ID of the current process
- gid_t getgid(void);
 - Returns the real group ID of the current process
- gid_t getegid(void);
 - Returns the effective group ID of the current process



Change UID and GID (1)

```
#include <unistd.h>
#include <sys/types.h>
int setuid( uid_t uid )
int setgid( gid t gid )
```

- Sets the effective user ID of the current process.
 - Superuser process resets the real effective user IDs to uid.
 - Non-superuser process can set effective user ID to uid, only when uid equals real user ID or the saved set-user ID (set by executing a setuid-program in exec).
 - In any other cases, setuid returns error.





Change UID and GID (2)

ID	exec		setuid(uid)	
	suid bit off	suid bit on	superuser	other users
real-uid effective-uid	unchanged unchanged	unchanged set from user ID of program	uid uid	unchanged uid
saved set-uid	copied from euid	file copied from euid	uid	unchanged



Change UID and GID (3)

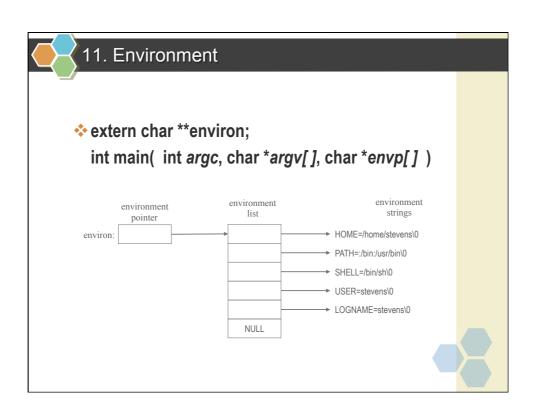
#include <unistd.h>
#include <sys/types.h>
int setreuid(uid t ruid, uid t euid)

- Sets real and effective user ID's of the current process.
- Un-privileged users may change the real user ID to the effective user ID and vice-versa.
- It is also possible to set the effective user ID from the saved user ID.
- Supplying a value of -1 for either the real or effective user ID forces the system to leave that ID unchanged.
- ❖ If the real user ID is changed or the effective user ID is set to a value not equal to the previous real user ID, the saved user ID will be set to the new effective user ID.



Change UID and GID (4)

- int seteuid(uid_t euid);
 - seteuid(euid) is functionally equivalent to setreuid(-1, euid).
 - Setuid-root program wishing to temporarily drop root privileges, assume the identity of a non-root user, and then regain root privileges afterwards cannot use setuid, because setuid issued by the superuser changes all three IDs. One can accomplish this with seteuid.
- int setregid(gid_t rgid, gid_t egid);
- int setegid(gid_t egid);



Example: environ

```
#include <stdio.h>

void main( int argc, char *argv[], char *envp[] )
   {
   int i;
   extern char **environ;

   printf( "from argument envp\n" );

   for( i = 0; envp[i]; i++ )
        puts( envp[i] );

   printf("\nFrom global variable environ\n");

   for( i = 0; environ[i]; i++ )
        puts(environ[i]);
}
```



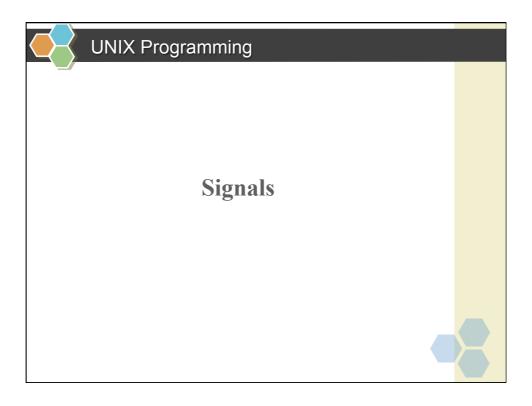
getenv

- #include <stdlib.h>
 - char *getenv(const char *name);
 - Searches the environment list for a string that matches the string pointed to by name.
 - Returns a pointer to the value in the environment, or NULL if there is no match.



- #include <stdlib.h>
 int putenv(const char *string);
 - Adds or changes the value of environment variables.
 - The argument string is of the form name=value.
 - If name does not already exist in the environment, then string is added to the environment.
 - If name does exist, then the value of name in the environment is changed to value.
 - Returns zero on success, or -1 if an error occurs.

#include <stdio.h> #include <stdib.h> void main(void) { printf("Home directory is %s\n", getenv("HOME")); putenv("HOME=/"); printf("New home directory is %s\n", getenv("HOME")); }



Overview

- 1. Definition
- 2. Signal Types
- 3. Generating a Signal
- 4. Responding to a Signal
- **5. Common Uses of Signals**
- 6. Timeout on a read()
- **7. POSIX Signal Functions**
- 8. Interrupted System Calls
- 9. System Calls inside Handlers
- 10. More Information





1. Definition

- *A signal is an asynchronous event which is delivered to a process.
- *Asynchronous means that the event can occur at any time
 - may be unrelated to the execution of the process
 - e.g. user types ctrl-C, or the modem hangs



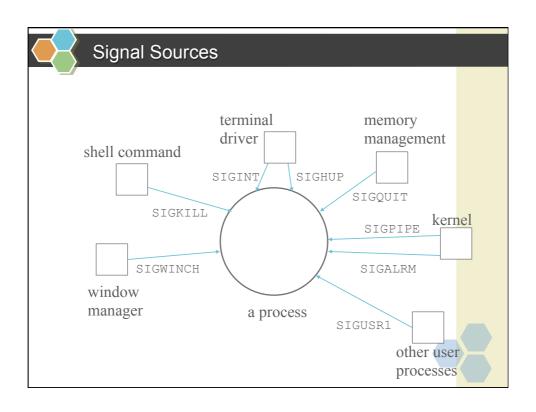


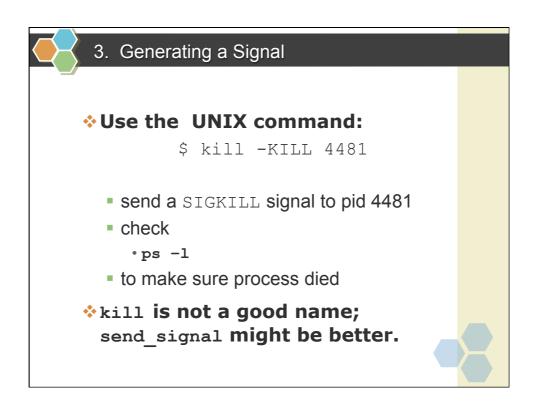
2. Signal Types (31 in POSIX)

*	<u>Name</u>	Description	Default Action	
	SIGINT	Interrupt character type	d terminate process	S
	SIGQUIT	Quit character typed (^\)	create core image	
	SIGKILL	kill -9	terminate process	
	SIGSEGV	Invalid memory reference	e create core image	е
	SIGPIPE	Write on pipe but no read	der terminate proce	ess
	SIGALRM	alarm() clock 'rings'		
	SIGUSR1	user-defined signal type		
	SIGUSR2	user-defined signal type	terminate process	

❖ See man 7 signal







kill()

- Send a signal to a process (or group of processes).
- *#include <signal.h>
 int kill(pid t pid, int signo);
- ❖ Return 0 if ok, -1 on error.





Some pid Values

❖ pid Meaning

> 0 send signal to process pid

== 0 send signal to all processes whose process group ID equals the sender's pgid.
e.g. parent kills all children





4. Responding to a Signal

- A process can:
 - ignore/discard the signal (not possible with SIGKILL Or SIGSTOP)
 - execute a signal handler function, and then possibly resume execution or terminate
 - carry out the default action for that signal
- The choice is called the process' signal disposition



signal(): library call

- Specify a signal handler function to deal with a signal type.
- #include <signal.h>
 typedef void Sigfunc(int); /* my defn */

Sigfunc *signal(int signo, Sigfunc *handler);

- signal returns a pointer to a function that returns an int (i.e. it returns a pointer to Sigfunc)
- *Returns previous signal disposition if ok, SIG_ERR on error.



Actual Prototype

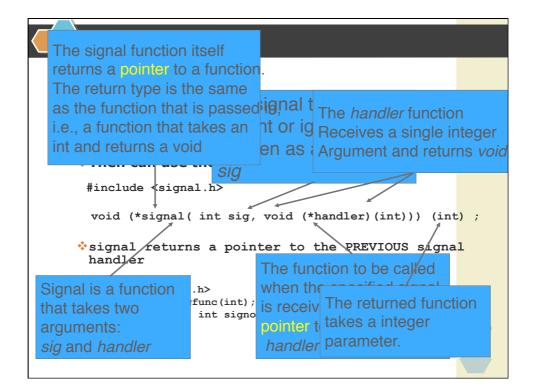
The actual prototype, listed in the "man" page is a bit perplexing but is an expansion of the Sigfune type:

```
void (*signal(int signo, void(*handler)(int))) (int);
```

♦In Linux:

```
typedef void (*sighandler_t) (int);
sig_handler_t signal(int signo, sighandler_t handler);
```

Signal returns a pointer to a function that returns an int



```
int main()
{
    signal( SIGINT, foo );
    :

    /* do usual things until SIGINT */
    return 0;
}

void foo( int signo )
{
    :     /* deal with SIGINT signal */
    return;    /* return to program */
}
```

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>

void sig_usr( int signo );  /* handles two signals */

int main()
{
  int i = 0;
  if( signal( SIGUSR1, sig_usr ) == SIG_ERR )
     printf( "Cannot catch SIGUSR1\n" );
  if( signal( SIGUSR2, sig_usr ) == SIG_ERR )
     printf("Cannot catch SIGUSR2\n");
  :

     continued
```

```
:
while(1)
{
    printf("%2d\n", I);
    pause();
    /* pause until signal handler
     * has processed signal */
    i++;
    }
    return 0;
}
```



Special Sigfunc * Values

⋄ Value	Meaning
sig_ign	Ignore / discard the signal.
SIG_DFL	Use default action to handle signal.
SIG_ERR	Returned by signal() as an error.





Multiple Signals

- ❖If many signals of the same type are waiting to be handled (e.g. two SIGINTS), then most UNIXs will only deliver one of them.
 - the others are thrown away
- *If many signals of different types are waiting to be handled (e.g. a SIGINT, SIGSEGV, SIGUSR1), they are not delivered in any fixed order.



pause()

- Suspend the calling process until a signal is caught.
- #include <unistd.h>
 int pause(void);
- ❖ Returns -1 with errno assigned EINTR. (Linux assigns it ERESTARTNOHAND).
- * pause () only returns after a signal handler has returned.





The Reset Problem

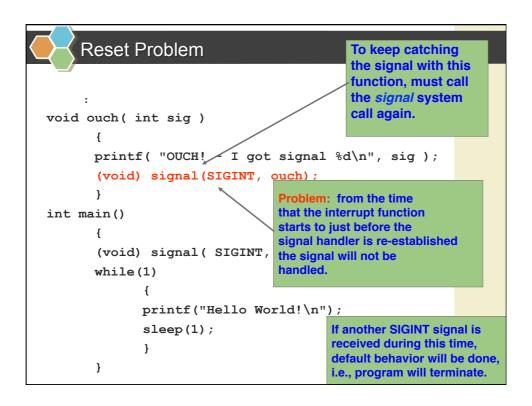
- In Linux (and many other UNIXs), the signal disposition in a process is reset to its default action immediately after the signal has been delivered.
- Must call signal() again to reinstall the signal handler function.



Reset Problem Example

```
int main()
    {
      signal(SIGINT, foo);
      :
      /* do usual things until SIGINT */
    }

void foo(int signo)
    {
      signal(SIGINT, foo); /* reinstall */
        :
      return;
    }
}
```





Re-installation may be too slow!

- *There is a (very) small time period in foo() when a new SIGINT signal will cause the default action to be carried out -- process termination.
- With signal() there is no answer to this problem.
 - POSIX signal functions solve it (and some other later UNIXs)



5. Common Uses of Signals

- 5.1. Ignore a Signal
- 5.2. Clean up and Terminate
- **5.3.** Dynamic Reconfiguration
- **5.4. Report Status**
- 5.5. Turn Debugging on/off
- **5.6.** Restore Previous Handler



5.1. Ignore a Signal

```
int main()
  signal(SIGINT, SIG IGN);
  signal(SIGQUIT, SIG IGN);
  /* do work without interruptions */
}
```

- **♦ Cannot ignore/handle** SIGKILL **or** SIGSTOP
- **♦ Should check for SIG ERR**





5.2. Clean up and Terminate

```
:
/* global variables */
int my_children_pids;
:
void clean_up(int signo);
int main()
{
   signal(SIGINT, clean_up);
   :
}
```

continued



Problems

- ❖If a program is run in the background then the interrupt and quit signals (SIGINT, SIGQUIT) are automatically ignored.
- Your code should not override these changes:
 - check if the signal dispositions are SIG IGN



Checking the Disposition

new disposition

old disposition

if(signal(SIGINT, SIG_IGN) != SIG_IGN)
 signal(SIGINT, clean_up);

if(signal(SIGQUIT, SIG_IGN) != SIG_IGN)
 signal(SIGQUIT, clean_up);
 :

Note: cannot check the signal disposition without changing it (sigaction that we will look at later, is different)



5.3. Dynamic Reconfiguration

```
:
void read_config(int signo);
int main()
{
  read_config(0); /* dummy argument */
  while (1)
    /* work forever */
}
```

```
void read_config(int signo)
{
   int fd;
   signal( SIGHUP, read_config );

   fd = open("config_file", O_RDONLY);
   /* read file and set global vars */
   close(fd);

   return;
}
```



Problems

- Reset problem
- Handler interruption
 - what is the effect of a SIGHUP in the middle of read config()'s execution?
- Can only affect global variables.



5.4. Report Status

```
void print_status(int signo)
{
    signal(SIGUSR1, print_status);
    printf("%d blocks copied\n", count);
    return;
}

* Reset problem
    count value not always defined.
    Must use global variables for status information
```

5.5. Turn Debugging on/off

```
:
void toggle_debug(int signo);
int debug = 0; /* initialize here */
int main()
{
    signal(SIGUSR2, toggle_debug);
    /* do work */
    if (debug == 1)
        printf("...");
    ...
}
```

```
void toggle_debug(int signo)
{
    signal(SIGUSR2, toggle_debug);
    debug = ((debug == 1) ? 0 : 1);
    return;
}
```

5.6. Restore Previous Handler

```
:
Sigfunc *old_hand;

/* set action for SIGTERM;
    save old handler */
old_hand = signal(SIGTERM, foobar);

/* do work */

/* restore old handler */
signal(SIGTERM, old_hand);
    :
```



6. Implementing a read() Timeout

- Put an upper limit on an operation that might block forever
 - **e.g.** read()
- **♦6.1.** alarm()
 - 6.2. Bad read() Timeout
 - **6.3.** setjmp() and longjmp()
 - 6.4. Better read() Timeout





6.1. alarm()

- Set an alarm timer that will 'ring' after a specified number of seconds
 - a SIGALRM signal is generated
- #include <unistd.h>
 long alarm(long secs);
- Returns 0 or number of seconds until previously set alarm would have 'rung'.





Some Tricky Aspects

- ❖A process can have at most one alarm timer running at once.
- ❖If alarm() is called when there is an existing alarm set then it returns the number of seconds remaining for the old alarm, and sets the timer to the new alarm value.
 - What do we do with the "old alarm value"?
- ❖An alarm(0) call causes the previous alarm to be cancelled.



6.2. Bad read() Timeout

```
#include <stdio.h>
#include <unistd.h>
#include <signal.h>
#define MAXLINE 512

void sig_alrm( int signo );
int main()
    {
    int n;
    char line[MAXLINE];
    :
```

continued

```
void sig_alrm(int signo)
/* do nothing, just handle signal */
{
   return;
}
```



Problems

- The code assumes that the read() call terminates with an error after being interrupted (talk about this later).
- ❖ Race Conditon: The kernel may take longer than 10 seconds to start the read() after the alarm() call.
 - the alarm may 'ring' before the read() starts
 - then the read() is not being timed; may block forever
 - Two ways two solve this:
 - setjmp
 - sigprocmask and sigsuspend



6.3. setjmp() and longjmp()

- ❖ In C we cannot use goto to jump to a label in another function
 - use setjmp() and longjmp() for those 'long jumps'
- Only uses which are good style:
 - error handling which requires a deeply nested function to recover to a higher level (e.g. back to main())
 - coding timeouts with signals





Prototypes

- *#include <setjmp.h>
 int setjmp(jmp_buf env);
- ❖ Returns 0 if called directly, non-zero if returning from a call to longjmp().
- *#include <setjmp.h>
 void longjmp(jmp buf env, int val);



Behavior

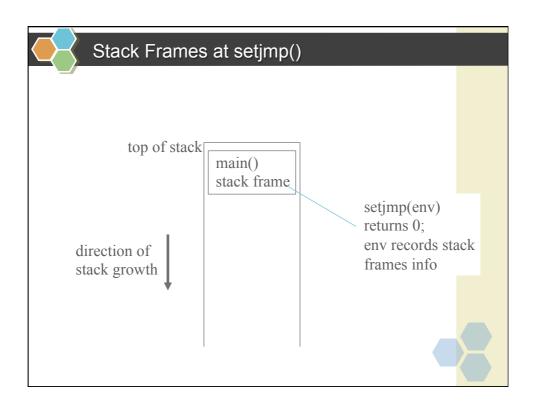
- ❖In the setjmp() call, env is initialized to information about the current state of the stack.
- ❖ The longjmp() call causes the stack to be reset to its env value.
- *Execution restarts after the
 setjmp() call, but this time setjmp()
 returns val.

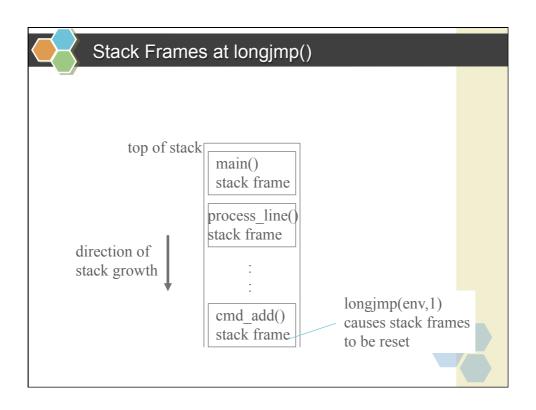
```
int main()
{
   char line[MAX];
   int errval;

   if(( errval = setjmp(env) ) != 0 )
        printf( "error %d: restart\n", errval );

   while( fgets( line, MAX, stdin ) != NULL )
        process_line(line);
   return 0;
}

continued
```





```
sleep1()
#include <signal.h>
#include <unistd.h>
void sig_alrm( int signo )
    return; /* return to wake up pause */
unsigned int sleep1( unsigned int nsecs )
    if( signal( SIGALRM, sig_alrm ) == SIG_ERR )
          return (nsecs);
                              /* starts timer */
    alarm( nsecs );
                             /* next caught signal wakes */
    pause();
                             /* turn off timer, return unslept
    return( alarm( 0 ) );
                               * time */
    }
```



Sleep1 and Sleep2

- Sleep2 fixes race condition. Even if the pause is never executed.
- There is one more problem (will talk about that after "fixing the earlier read function")





Status of Variables?

- ❖ The POSIX standard says:
 - global and static variable values will not be changed by the longjmp() call
- Nothing is specified about local variables, are they "rolled back" to their original values (at the setjmp call) as the stack"?
 - they may be restored to their values at the first setjmp(), but maybe not
 - Most implementations do not roll back their values



#include <stdio.h> #include <unistd.h> #include <setjmp.h> #include <signal.h> #define MAXLINE 512 void sig_alrm(int signo); jmp_buf env_alrm; int main() { int n; char line[MAXLINE]; : continued

```
void sig_alrm(int signo)
/* interrupt the read() and jump to
    setjmp() call with value 1
*/
{
    longjmp(env_alrm, 1);
}
```



Caveat: Non-local Jumps

From the UNIX man pages:

WARNINGS

If longjmp() or siglongjmp() are called even though env was never primed by a call to setjmp() or sigsetjmp(), or when the last such call was in a function that has since returned, absolute chaos is guaranteed.





A Problem Remains!

If the program has several signal handlers then:

- execution might be inside one when an alarm 'rings'
- the longjmp() call will jump to the setjmp() location, and abort the other signal handler -- might lose / corrupt data





7. POSIX Signal Functions

- The POSIX signal functions can control signals in more ways:
 - can block signals for a while, and deliver them later (good for coding critical sections)
 - can switch off the resetting of the signal disposition when a handler is called (no reset problem)



The POSIX signal system, uses signal sets, to deal with pending signals that might otherwise be missed while a signal is being processed





7.1. Signal Sets

- The signal set stores collections of signal types.
- Sets are used by signal functions to define which signal types are to be processed.
- POSIX contains several functions for creating, changing and examining signal sets.



Prototypes

#include <signal.h>





7.2. sigprocmask()

- ❖ A process uses a signal set to create a mask which defines the signals it is blocking from delivery. good for critical sections where you want to block certain signals.
- how indicates how mask is modified



how Meanings

❖ Value Meaning

SIG_BLOCK set signals are added to mask

SIG_UNBLOCK set signals are removed from mask

SIG_SETMASK set becomes new mask



A Critical Code Region

```
sigset_t newmask, oldmask;
sigemptyset( &newmask );
sigaddset( &newmask, SIGINT );

/* block SIGINT; save old mask */
sigprocmask( SIG_BLOCK, &newmask, &oldmask );

/* critical region of code */

/* reset mask which unblocks SIGINT */
sigprocmask( SIG_SETMASK, &oldmask, NULL );
```

7.3. sigaction()

- \$Supercedes (more powerful than)
 signal()
 - sigaction() can be used to code a nonresetting signal()
- #include <signal.h>



sigaction Structure

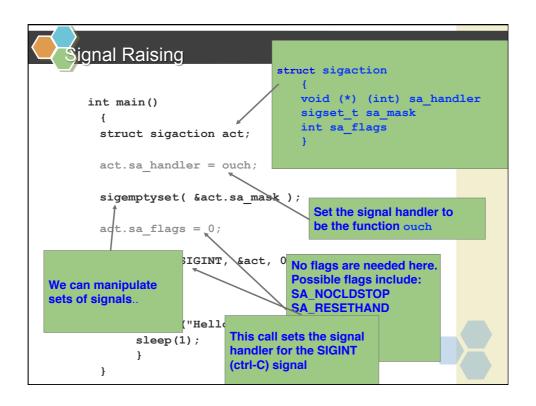
sa_flags -

- SIG_DFL reset handler to default upon return
- SA_SIGINFO denotes extra information is passed to handler (.i.e. specifies the use of the "second" handler in the structure.



sigaction() Behavior

- A signo signal causes the sa_handler signal handler to be called.
- While sa_handler executes, the signals in sa_mask are blocked. Any more signals are also blocked.
- * sa_handler remains installed until it is changed by another sigaction() call. No reset problem.





Signal Raising

- This function will continually capture the ctrl-C (SIGINT) signal.
- Default behavior is not restored after signal is caught.
- ❖To terminate the program, must type ctrl-\, the SIGQUIT signal.



sigexPOS.c

```
/* sigexPOS.c - demonstrate sigaction() */
/* include files as before */
int main(void)
{
    /* struct to deal with action on signal set */
    static struct sigaction act;

    void catchint(int); /* user signal handler */
    /* set up action to take on receipt of SIGINT */
    act.sa_handler = catchint;
```

```
/* create full set of signals */
sigfillset(&(act.sa_mask));

/* before sigaction call, SIGINT will terminate
  * process */

/* now, SIGINT will cause catchint to be executed */
sigaction( SIGINT, &act, NULL );
sigaction( SIGQUIT, &act, NULL );

printf("sleep call #1\n");
sleep(1);

/* rest of program as before */
```



Signals - Ignoring signals

- Other than SIGKILL and SIGSTOP, signals can be ignored:
- Instead of in the previous program:

```
act.sa_handler = catchint /* or whatever */
We use:
   act.sa_handler = SIG_IGN;
The ^C key will be ignored
```



Restoring previous action

The third parameter to sigaction, oact, can be used:

```
/* save old action */
sigaction( SIGTERM, NULL, &oact );

/* set new action */
act.sa_handler = SIG_IGN;

sigaction( SIGTERM, &act, NULL );

/* restore old action */
sigaction( SIGTERM, &oact, NULL );
```



A Basic signal()



7.4. Other POSIX Functions

sigpending() examine blocked signals

sigsetjmp()
siglongjmp()

jump functions for use in signal handlers which handle masks correctly

sigsuspend()
atomically reset mask
and sleep





[sig]longjmp & [sig]setjmp

NOTES (longjmp, sigjmp)

POSIX does not specify whether longjmp will restore the signal context. If you want to save and restore signal masks, use siglongjmp.

NOTES (setjmp, sigjmp)

POSIX does not specify whether setjmp will save the signal context. (In SYSV it will not. In BSD4.3 it will, and there is a function _setjmp that will not.) If you want to save signal masks, use sigsetjmp.



```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>

sigjmp_buf buf;

void handler(int sig)
    {
        siglongjmp(buf, 1);
    }

main()
{
    signal(SIGINT, handler);

    if( sigsetjmp(buf, 1) == 0 )
        printf("starting\n");
    else
        printf("restarting\n");
...
```

```
while(1)
   {
       sleep(5);
       printf(" waiting...\n");
}
> a.out
starting
 waiting...
 waiting... ← Control-c
restarting
 waiting...
  waiting...
 waiting... Control-c
restarting
 waiting... ←
                  Control-c
restarting
  waiting...
  waiting...
```



8. Interrupted System Calls

- When a system call (e.g. read()) is interrupted by a signal, a signal handler is called, returns, and then what?
- On many UNIXs, slow system function calls do not resume. Instead they return an error and errno is assigned EINTR.
 - true of Linux, but can be altered with (Linuxspecific) siginterrupt()



Slow System Functions

- Slow system functions carry out I/O on things that can possibly block the caller forever:
 - pipes, terminal drivers, networks
 - some IPC functions
 - pause(), some uses of ioctl()
- Can use signals on slow system functions to code up timeouts (e.g. did earlier)





Non-slow System Functions

- Most system functions are non-slow, including ones that do disk I/O
 - e.g. read() of a disk file
 - read () is sometimes a slow function, sometimes not
- Some UNIXs resume non-slow system functions after the handler has finished.
- Some UNIXs only call the handler after the non-slow system function call has finished.



9. System Calls inside Handlers

- If a system function is called inside a signal handler then it may interact with an interrupted call to the same function in the main code.
 - e.g. malloc()
- This is not a problem if the function is reentrant
 - a process can contain multiple calls to these functions at the same time
 - e.g. read(), write(), fork(), many more





Non-reentrant Functions

- ❖A functions may be non-reentrant (only one call to it at once) for a number of reasons:
 - it uses a static data structure
 - it manipulates the heap: malloc(), free(), etc.
 - it uses the standard I/O library
 - e,g, scanf(), printf()
 - the library uses global data structures in a non-reentrant way





errno Problem

- *errno is usually represented by a global variable.
- Its value in the program can be changed suddenly by a signal handler which produces a new system function error.

