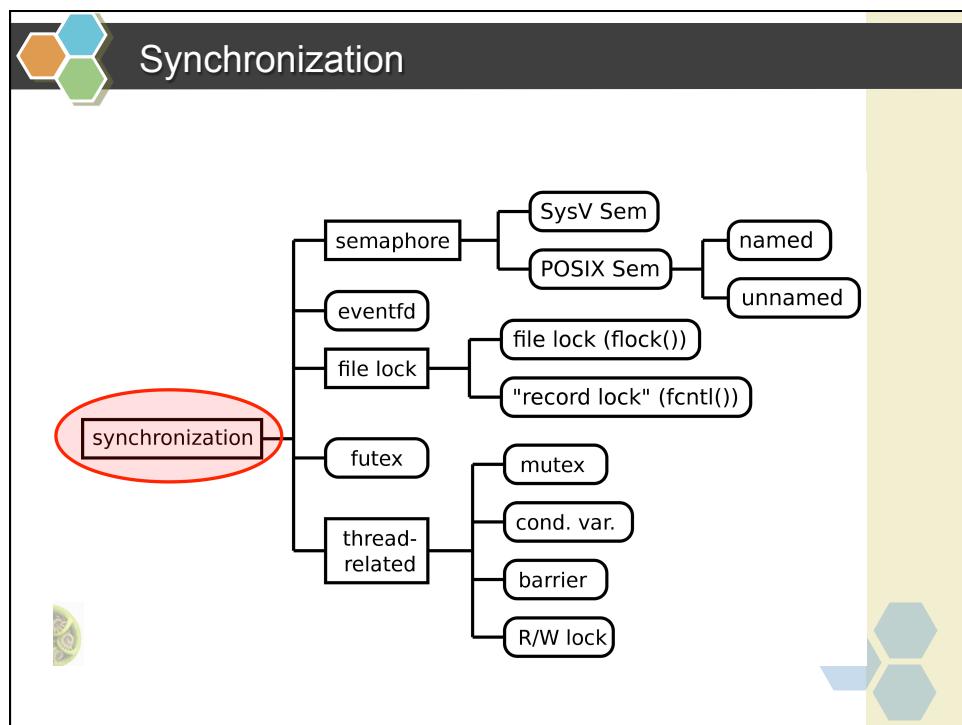
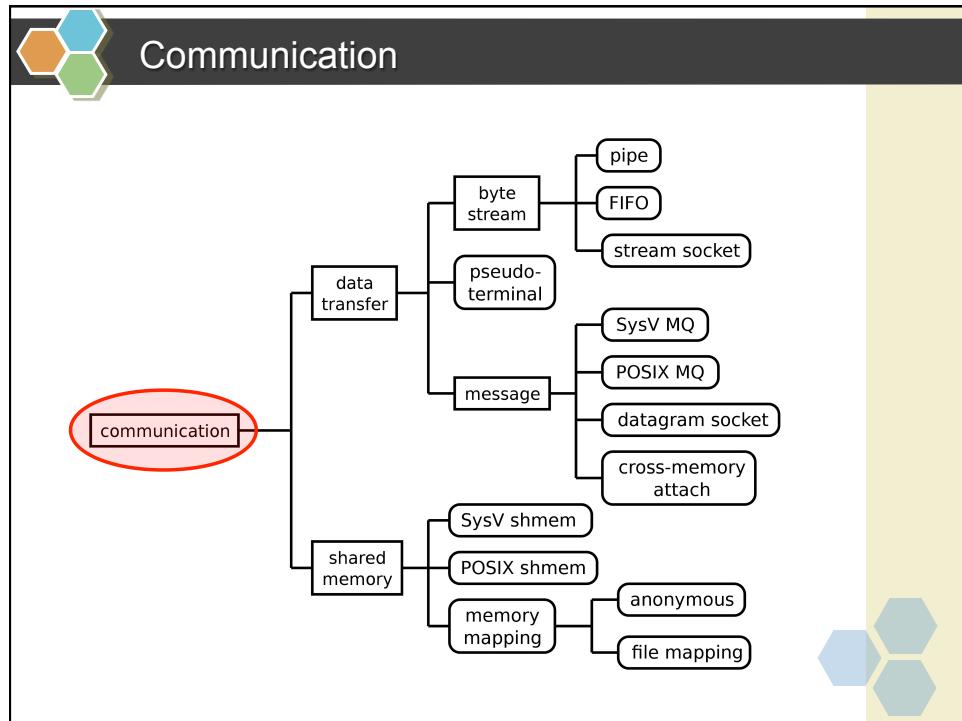




A lot of IPC

- ❖ Pipes
- ❖ FIFOs
- ❖ Sockets
 - Stream vs Datagram (vs Seq. packet)
 - UNIX vs Internet domain
- ❖ POSIX message queues
- ❖ POSIX shared memory
- ❖ POSIX semaphores
 - Named, Unnamed
- ❖ Shared memory mappings
 - File vs Anonymous
- ❖ Cross-memory attach
 - `proc_vm_readv()` / `proc_vm_writev()`
- ❖ Signals
 - Standard, Realtime
- ❖ File locks
- ❖ Mutexes
- ❖ Read-write locks



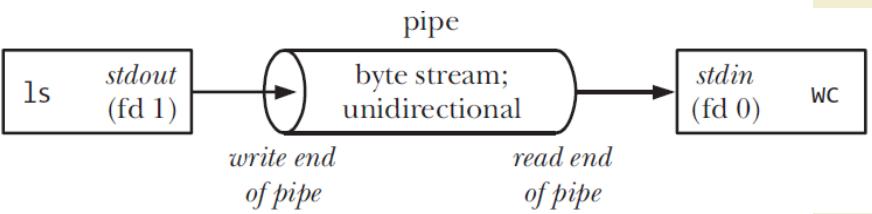
 Outline

- ❖ **Communication Techniques**
 - **Pipes**
 - FIFO (named pipes)
 - POSIX Message Queues
 - Shared Memory
 - Shared Anonymous Mapping
 - Shared File Mapping
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores



 Pipes

- ❖ **ls | wc -l**



```

graph LR
    A["ls  
stdout  
(fd 1)"] --> B("byte stream;  
unidirectional")
    B --> C["stdin  
(fd 0)  
wc"]
    style B fill:none,stroke:none
    
```

The diagram illustrates a pipe connection between two processes: 'ls' and 'wc'. The 'ls' process is shown on the left with its standard output (stdout, fd 1) connected to the write end of a pipe. The 'wc' process is on the right with its standard input (stdin, fd 0) connected to the read end of the same pipe. The pipe is represented by a horizontal oval labeled 'byte stream; unidirectional'. The text 'write end of pipe' is positioned below the arrow from 'ls' to the pipe, and 'read end of pipe' is positioned below the arrow from the pipe to 'wc'.



Pipes

- ❖ Pipe == byte stream buffer in kernel
 - Sequential (can't lseek())
 - Multiple readers/writers difficult
- ❖ Unidirectional
 - Write end + read end

Creating and using pipe

- ❖ **Created using `pipe()`:**

```

int filedes[2];
pipe(filedes);
...
write(filedes[1], buf, count);
read(filedes[0], buf, count);

```

The diagram illustrates a pipe mechanism. On the left, a rounded rectangle labeled "calling process" contains the variable "filedes[1]" above "filedes[0]". A curved arrow points from this rectangle to a horizontal oval labeled "pipe". Inside the "pipe" oval, there is a horizontal arrow pointing to the right, labeled "direction of data flow".

Sharing a pipe

- ❖ **Pipes are anonymous**
 - No name in file system
- ❖ **How do two processes share a pipe?**

Sharing a pipe

```
int filedes[2];
pipe(filedes);
child_pid = fork();
```

- `fork()` duplicates parent's file descriptors

The diagram illustrates the sharing of a pipe between a parent process and a child process. It shows two rectangular boxes representing the processes. The top box is labeled "parent process" and contains "filedes[1]" and "filedes[0]". The bottom box is labeled "child process" and also contains "filedes[1]" and "filedes[0]". A horizontal pipe symbol, consisting of two concentric ovals with an arrow pointing from left to right, connects the two boxes. Arrows point from each "filedes[0]" descriptor to the left side of the pipe, and from each "filedes[1]" descriptor to the right side of the pipe, indicating that both processes can read from the parent's write end and write to the parent's read end.

Sharing a pipe

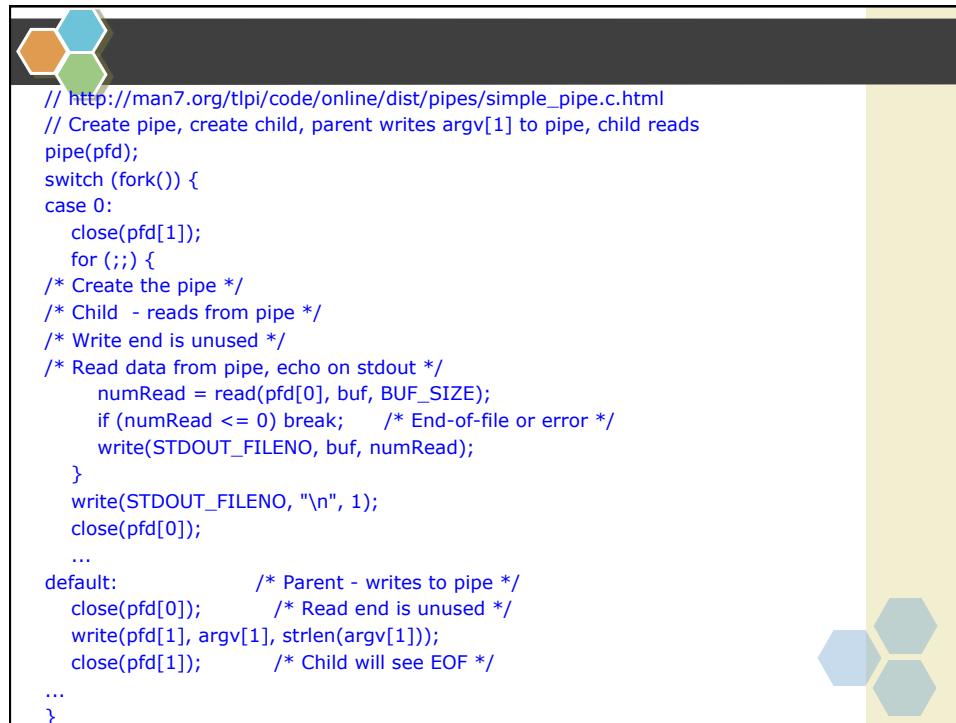
```
int filedes[2];
pipe(filedes);
child_pid = fork();
if (child_pid == 0) {
    close(filedes[1]);
    /* Child now reads */
} else {
    close(filedes[0]);
    /* Parent now writes */
}
(error checking omitted!)
```

The diagram shows a vertical stack of three hexagons (orange, blue, green) on the left and a vertical bar of three hexagons (blue, grey, light blue) on the right. In the center, there is a diagram of a pipe. A box labeled "parent process" contains "filedes[1]". A box labeled "child process" contains "filedes[0]". A horizontal pipe connects them. Arrows point from the parent's writing end to the child's reading end.

Closing unused file descriptors

- ❖ Parent and child must close unused descriptors
 - Necessary for correct use of pipes!
- ❖ close() write end
 - read() returns 0 (EOF)
- ❖ close() read end
 - write() fails with EPIPE error + SIGPIPE signal

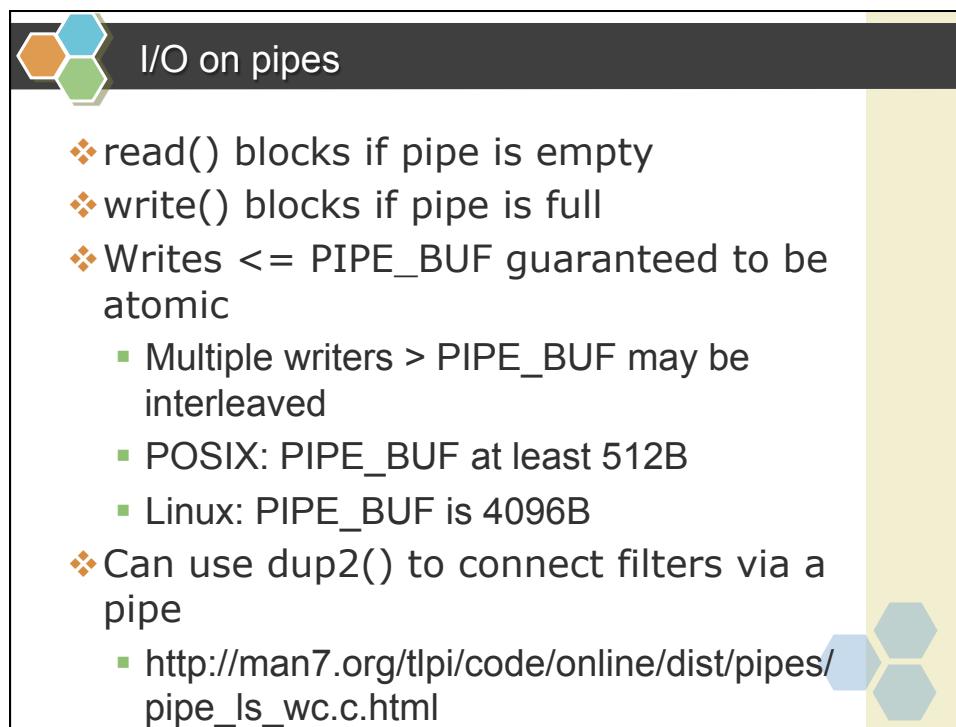
The diagram shows a vertical stack of three hexagons (orange, blue, green) on the left and a vertical bar of three hexagons (blue, grey, light blue) on the right.



```

// http://man7.org/tlpi/code/online/dist/pipes/simple_pipe.c.html
// Create pipe, create child, parent writes argv[1] to pipe, child reads
pipe(pfd);
switch (fork()) {
case 0:
    close(pfd[1]);
    for (;;) {
        /* Create the pipe */
        /* Child - reads from pipe */
        /* Write end is unused */
        /* Read data from pipe, echo on stdout */
        numRead = read(pfd[0], buf, BUF_SIZE);
        if (numRead <= 0) break; /* End-of-file or error */
        write(STDOUT_FILENO, buf, numRead);
    }
    write(STDOUT_FILENO, "\n", 1);
    close(pfd[0]);
    ...
default:           /* Parent - writes to pipe */
    close(pfd[0]);      /* Read end is unused */
    write(pfd[1], argv[1], strlen(argv[1]));
    close(pfd[1]);      /* Child will see EOF */
    ...
}

```

- 
- ## I/O on pipes
- ❖ `read()` blocks if pipe is empty
 - ❖ `write()` blocks if pipe is full
 - ❖ Writes \leq PIPE_BUF guaranteed to be atomic
 - Multiple writers > PIPE_BUF may be interleaved
 - POSIX: PIPE_BUF at least 512B
 - Linux: PIPE_BUF is 4096B
 - ❖ Can use `dup2()` to connect filters via a pipe
 - http://man7.org/tlpi/code/online/dist/pipes/pipe_ls_wc.c.html

 Pipes have limited capacity

- ❖ **Limited capacity**
 - If pipe fills, write() blocks
 - Before Linux 2.6.11: 4096 bytes
 - Since Linux 2.6.11: 65,536 bytes
- ❖ **Apps should be designed not to care about capacity**
 - But, Linux has fcntl(fd, F_SETPIPE_SZ, size)
 - (not portable)



 Outline

- ❖ **Communication Techniques**
 - Pipes
 - **FIFO (named pipes)**
 - POSIX Message Queues
 - Shared Memory
 - Shared Anonymous Mapping
 - Shared File Mapping
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores



FIFO (named pipe)

- ❖ FIFO (named pipe)
- ❖ (Anonymous) pipes can only be used by related processes
- ❖ FIFOs == pipe with name in file system
Creation:
 - `mkfifo(pathname, permissions)`
- ❖ Any process can open and use FIFO I/O is same as for pipes

Opening a FIFO

- ❖ `open(pathname, O_RDONLY)`
 - Open read end
- ❖ `open(pathname, O_WRONLY)`
 - Open write end
- ❖ `open()` locks until other end is opened
 - Opens are synchronized
 - `open(pathname, O_RDONLY | O_NONBLOCK)` can be useful

 Outline

- ❖ **Communication Techniques**
 - Pipes
 - FIFO (named pipes)
 - **POSIX Message Queues**
 - Shared Memory
 - Shared Anonymous Mapping
 - Shared File Mapping
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores



 Highlights of POSIX MQs

- ❖ **Message-oriented communication**
 - Receiver reads messages one at a time
 - No partial or multiple message reads
 - Unlike pipes, multiple readers/writers can be useful
- ❖ **Messages have priorities**
 - Delivered in priority order
- ❖ **Message notification feature**



 **POSIX MQ API**

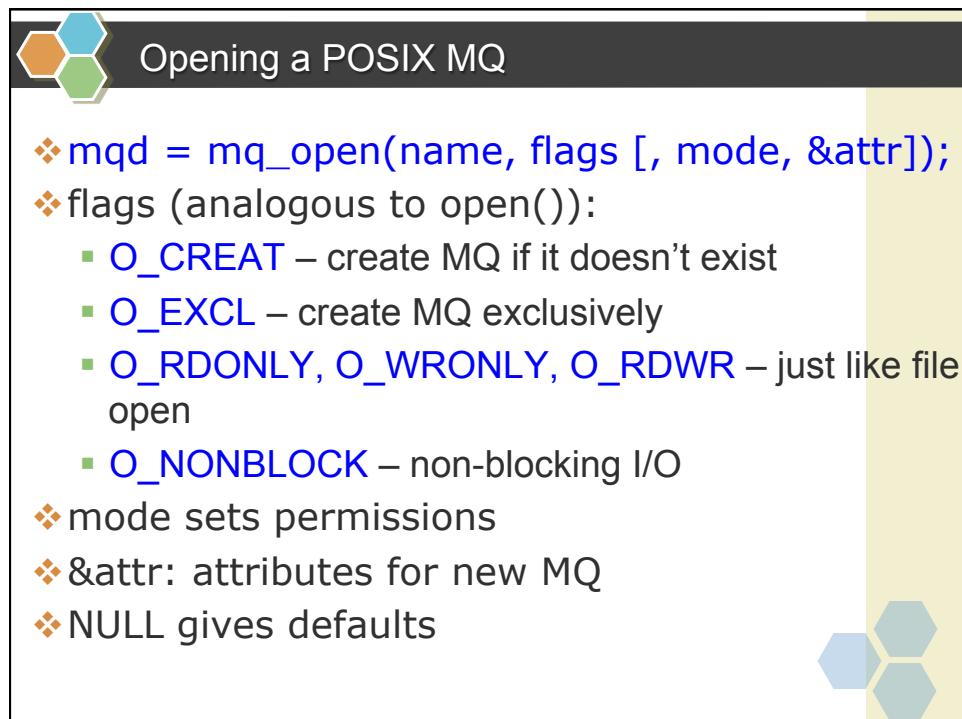
- ❖ **Queue management (analogous to files)**
 - `mq_open()`: open/create MQ, set attributes
 - `mq_close()`: close MQ
 - `mq_unlink()`: remove MQ pathname
- ❖ **I/O:**
 - `mq_send()`: send message
 - `mq_receive()`: receive message
- ❖ **Other:**
 - `mq_setattr()`, `mq_getattr()`: set/get MQ attributes
 - `mq_notify()`: request notification of msg arrival



 **Opening a POSIX MQ**

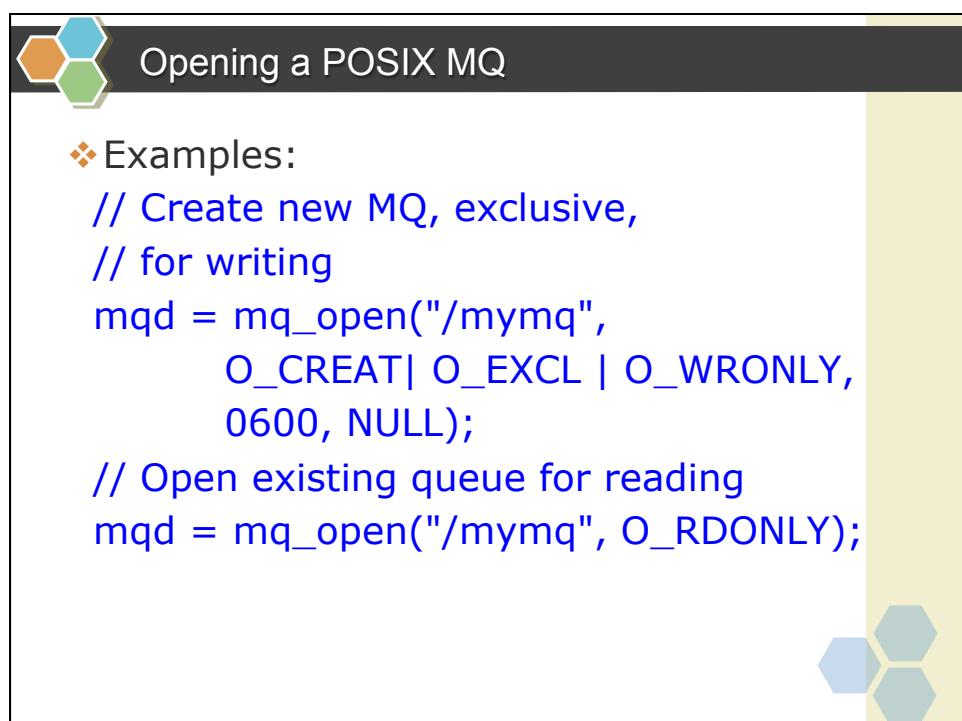
- ❖ `mqd = mq_open(name, flags [, mode, &attr]);`
- ❖ Open+create new MQ / open existing MQ
name has form /somename
 - Visible in a pseudo-filesystem
- ❖ Returns `mqd_t`, a message queue descriptor
 - Used by rest of API





Opening a POSIX MQ

- ❖ `mqd = mq_open(name, flags [, mode, &attr]);`
- ❖ flags (analogous to `open()`):
 - `O_CREAT` – create MQ if it doesn't exist
 - `O_EXCL` – create MQ exclusively
 - `O_RDONLY`, `O_WRONLY`, `O_RDWR` – just like file open
 - `O_NONBLOCK` – non-blocking I/O
- ❖ mode sets permissions
- ❖ &attr: attributes for new MQ
- ❖ NULL gives defaults



Opening a POSIX MQ

- ❖ Examples:

```
// Create new MQ, exclusive,  
// for writing  
mqd = mq_open("/mymq",  
              O_CREAT| O_EXCL | O_WRONLY,  
              0600, NULL);  
  
// Open existing queue for reading  
mqd = mq_open("/mymq", O_RDONLY);
```



Unlink a POSIX MQ

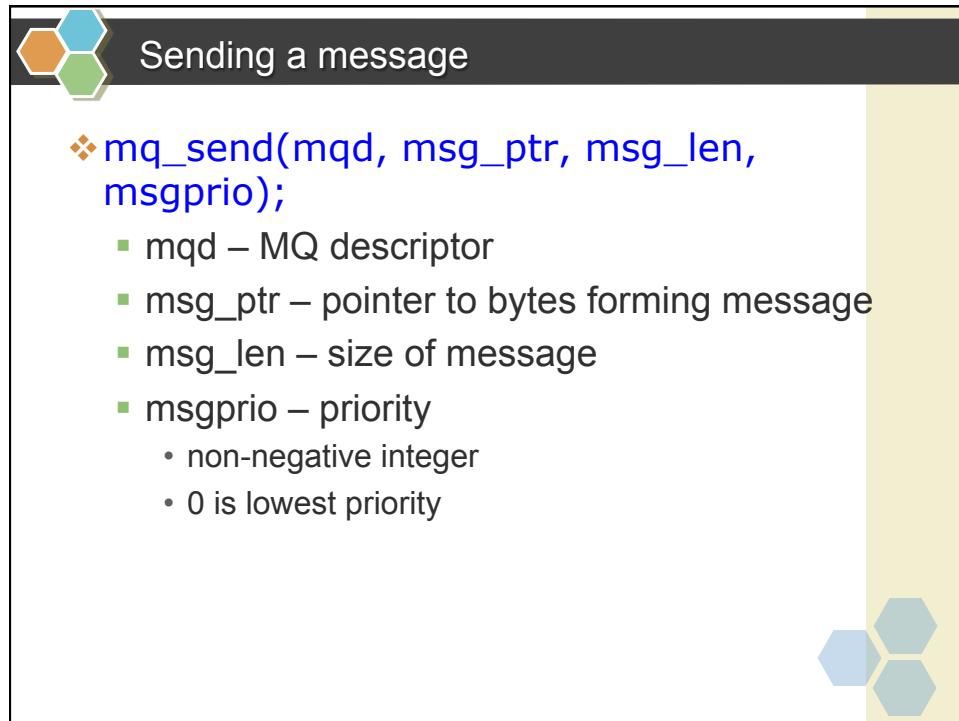
- ❖ **mq_unlink(name);**
- ❖ **MQs are reference-counted**
 - ==> MQ removed only after all users have closed



Nonblocking I/O on POSIX MQs

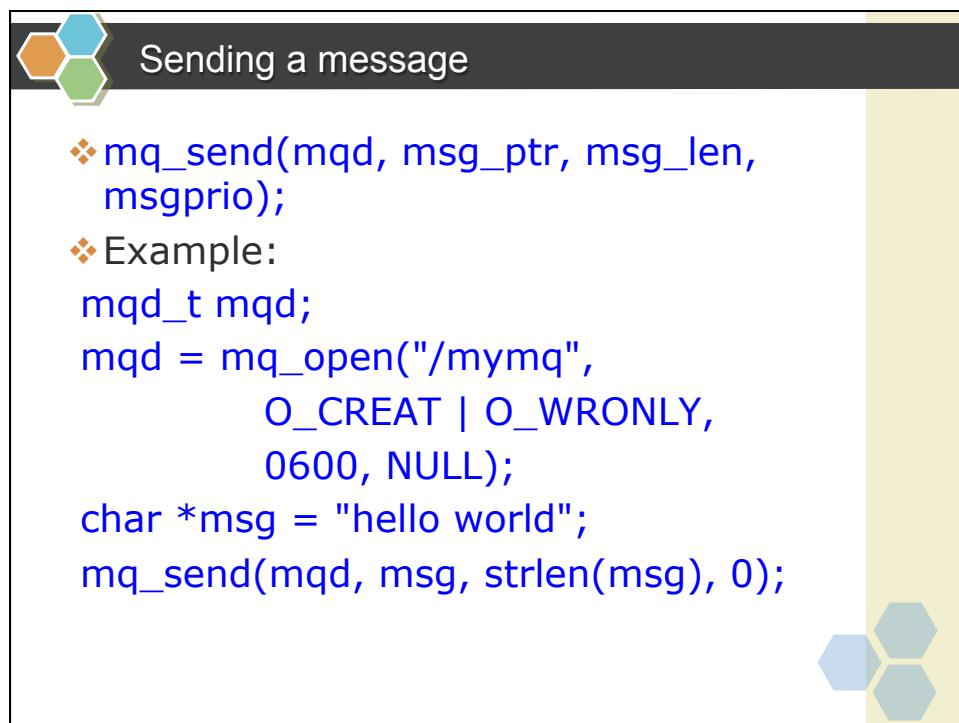
- ❖ **Message ques have a limited capacity**
 - Controlled by attributes
- ❖ **By default:**
 - `mq_receive()` blocks if no messages in queue
 - `mq_send()` blocks if queue is full
- ❖ **O_NONBLOCK:**
 - EAGAIN error instead of blocking
 - Useful for emptying queue without blocking





Sending a message

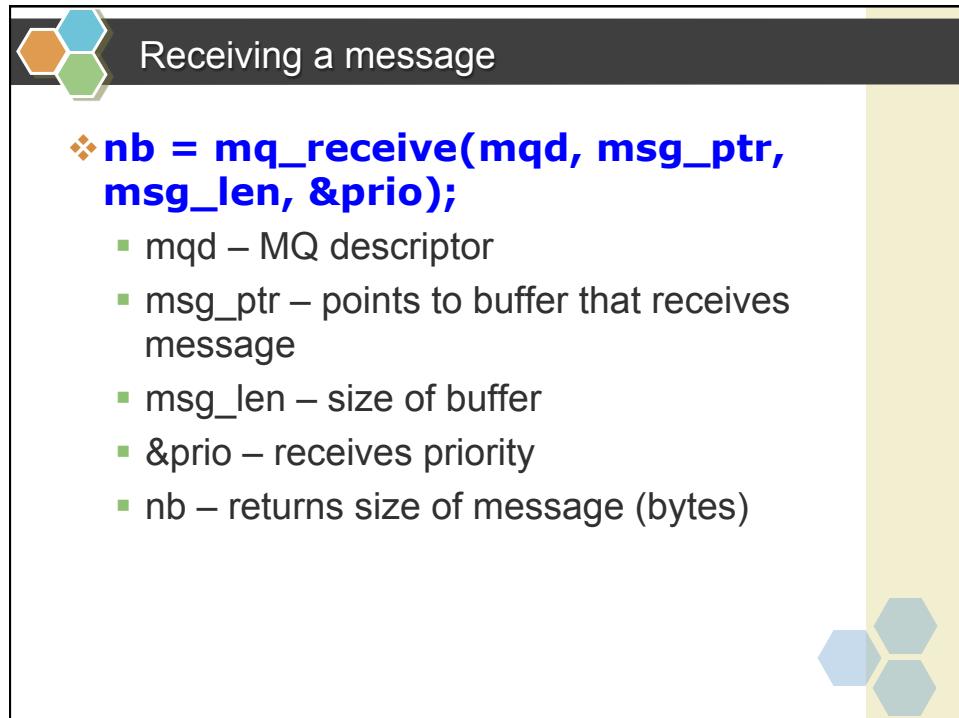
- ❖ `mq_send(mqd, msg_ptr, msg_len, msgprio);`
 - `mqd` – MQ descriptor
 - `msg_ptr` – pointer to bytes forming message
 - `msg_len` – size of message
 - `msgprio` – priority
 - non-negative integer
 - 0 is lowest priority



Sending a message

- ❖ `mq_send(mqd, msg_ptr, msg_len, msgprio);`
- ❖ Example:

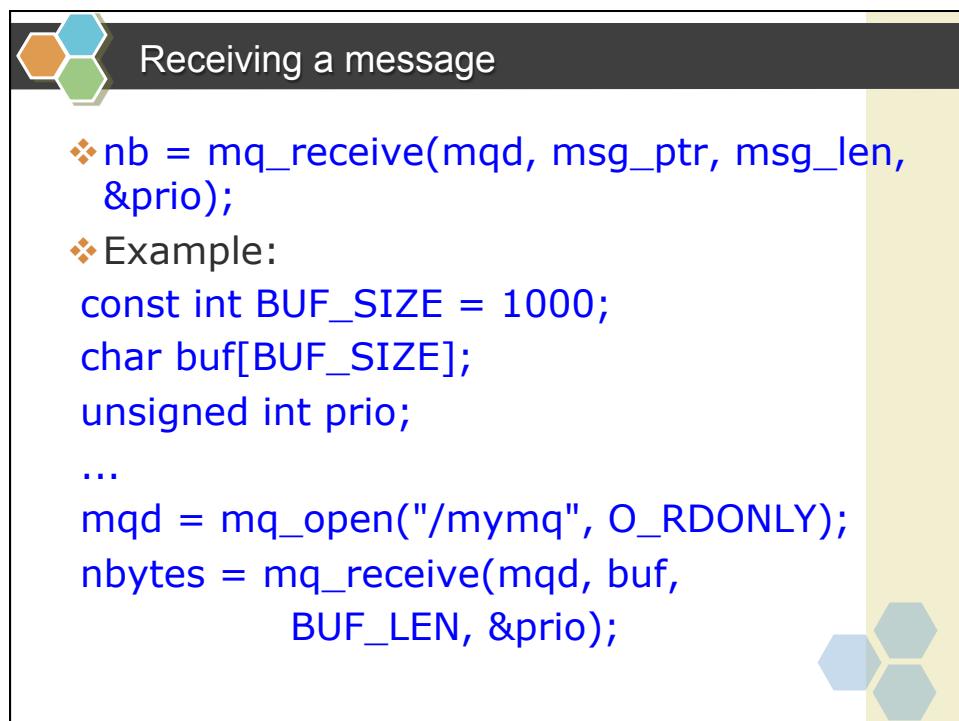
```
mqd_t mqd;
mqd = mq_open("/mymq",
              O_CREAT | O_WRONLY,
              0600, NULL);
char *msg = "hello world";
mq_send(mqd, msg, strlen(msg), 0);
```



Receiving a message

❖ **nb = mq_receive(mqd, msg_ptr, msg_len, &prio);**

- mqd – MQ descriptor
- msg_ptr – points to buffer that receives message
- msg_len – size of buffer
- &prio – receives priority
- nb – returns size of message (bytes)



Receiving a message

❖ **nb = mq_receive(mqd, msg_ptr, msg_len, &prio);**

❖ Example:

```
const int BUF_SIZE = 1000;
char buf[BUF_SIZE];
unsigned int prio;
...
mqd = mq_open("/mymq", O_RDONLY);
nbytes = mq_receive(mqd, buf,
                    BUF_LEN, &prio);
```

 **POSIX MQ notifications**

- ❖ `mq_notify(mqd, notification);`
- ❖ One process can register to receive notification Notified when new msg arrives on empty queue
- ❖ & only if another process is not doing `mq_receive()` notification says how caller should be notified
- ❖ Send me a signal
- ❖ Start a new thread (see `mq_notify(3)` for example)
- ❖ One-shot; must re-enable
- ❖ Do so before emptying queue!



 **POSIX MQ attributes**

```
struct mq_attr {
    long mq_flags;      // MQ description flags
                        // 0 or O_NONBLOCK
                        // [mq_getattr(), mq_setattr()]
    long mq_maxmsg;    // Max. # of msgs on queue
                        // [mq_open(), mq_getattr()]
    long mq_msgsize;   // Max. msg size (bytes)
                        // [mq_open(), mq_getattr()]
    long mq_curmsgs;  // # of msgs currently in queue
                        // [mq_getattr()]
};
```



 **POSIX MQ details**

- ❖ Per-process and system-wide limits govern resource usage
- ❖ Can mount filesystem to obtain info on MQs:


```
# mkdir /dev/mqueue
# mount -t mqueue none /dev/mqueue
# ls /dev/mqueue
mymq
# cat /dev/mqueue/mymq
QSIZE:129 NOTIFY:2 SIGNO:0 NOTIFY_PID:
8260
```



 **Outline**

- ❖ **Communication Techniques**
 - Pipes
 - FIFO (named pipes)
 - POSIX Message Queues
 - **Shared Memory**
 - Shared Anonymous Mapping
 - Shared File Mapping
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores



Shared memory

- ❖ Processes share same physical pages of memory
- ❖ Communication == copy data to memory
- ❖ Efficient; compare
 - Data transfer: user space ==> kernel ==> user space
 - Shared memory: single copy in user space
- ❖ But, need to synchronize access...

Shared memory

- ❖ Processes share physical pages of memory

The diagram illustrates the mapping of physical pages between two processes and physical memory. It shows two page tables (Process A page table and Process B page table) and a central column of physical pages.

- Process A page table:** Contains a header and "PT entries for mapped region".
- Process B page table:** Contains a header and "PT entries for mapped region".
- Physical memory:** Contains a header and "Mapped pages".
- Mapping:** Arrows from both page tables point to specific pages in the "Mapped pages" column. Specifically, the first entry in each page table points to the first page in the "Mapped pages" column.



Shared memory

- ❖ **We'll cover three types:**
 - Shared anonymous mappings
 - related processes
 - Shared file mappings
 - unrelated processes, backed by file in traditional filesystem
- ❖ **POSIX shared memory**
 - unrelated processes, without use of traditional filesystem



mmap()

- ❖ Syscall used in all three shmem types
- ❖ Rather complex:

```
void *mmap(void *daddr, size_t len, int prot, int flags, int fd, off_t offset);
```



mmap()

- ❖ `addr = mmap(daddr, len, prot, flags, fd, offset);`
- ❖ **daddr** – choose where to place mapping;
 - Best to use NULL, to let kernel choose len – size of mapping
- ❖ **prot** – memory protections (read, write, exec) flags – control behavior of call
 - MAP_SHARED, MAP_ANONYMOUS
- ❖ **fd** – file descriptor for file mappings
- ❖ **offset** – starting offset for mapping from file addr – returns address used for mapping

Using shared memory

- ❖ `addr = mmap(daddr, len, prot, flags, fd, offset);`
- ❖ **addr** looks just like any C pointer
- ❖ But, changes to region seen by all process that map it

The diagram illustrates the concept of shared memory using the `mmap()` function. It shows a vertical stack of memory regions labeled "Process virtual memory". A specific region within this stack is highlighted and labeled "mapped region". An arrow points upwards from the text "increasing memory address" to the top of the "mapped region". Another arrow points to the right from the text "address returned by `mmap()`" to the left boundary of the "mapped region".

 Outline

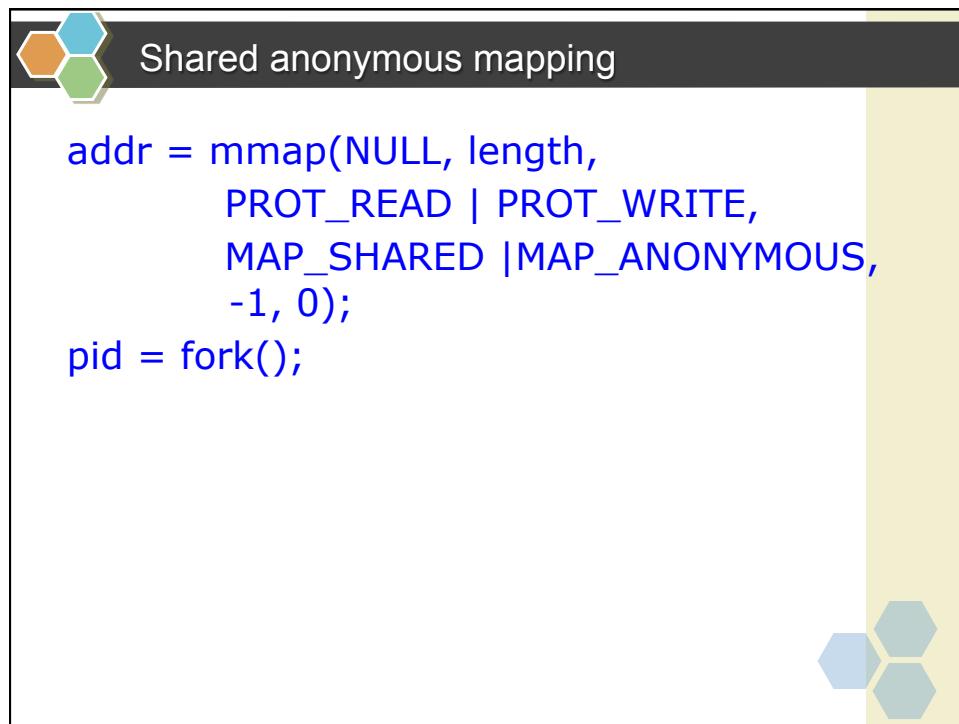
- ❖ **Communication Techniques**
 - Pipes
 - FIFO (named pipes)
 - POSIX Message Queues
 - Shared Memory
 - **Shared Anonymous Mapping**
 - Shared File Mapping
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores



 Shared anonymous mapping

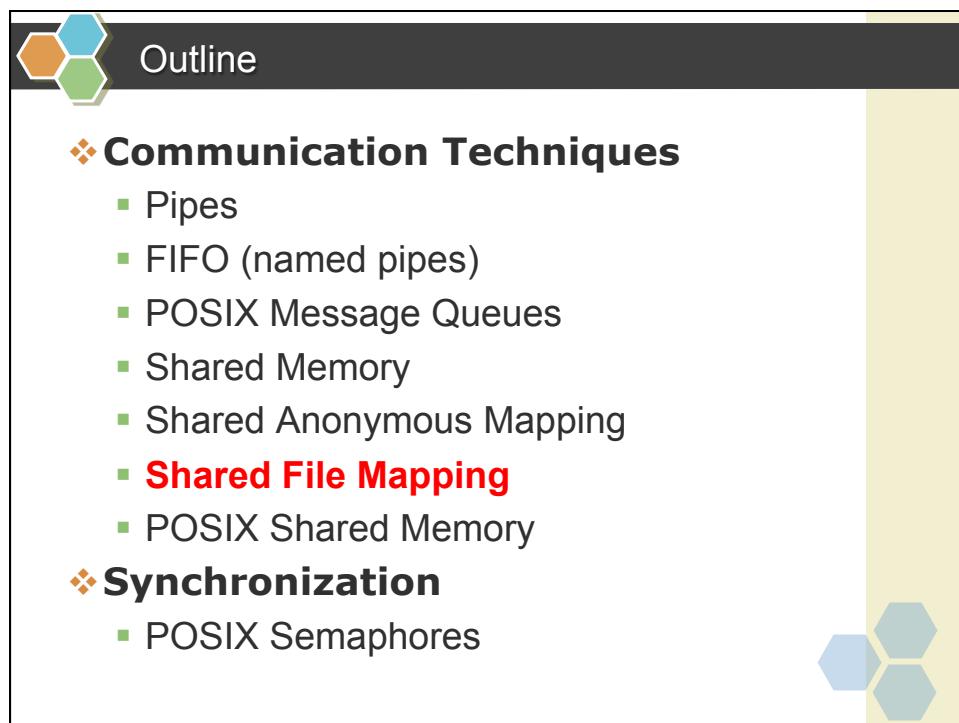
- ❖ Share memory between related processes
- ❖ `mmap()` fd and offset args unneeded
`addr = mmap(NULL, length,`
`PROT_READ | PROT_WRITE,`
`MAP_SHARED | MAP_ANONYMOUS, -1, 0);`
`pid = fork();`
- ❖ Allocates zero-initialized block of length bytes
- ❖ Parent and child share memory at
`addr:length`





Shared anonymous mapping

```
addr = mmap(NULL, length,  
           PROT_READ | PROT_WRITE,  
           MAP_SHARED | MAP_ANONYMOUS,  
           -1, 0);  
pid = fork();
```



Outline

- ❖ **Communication Techniques**
 - Pipes
 - FIFO (named pipes)
 - POSIX Message Queues
 - Shared Memory
 - Shared Anonymous Mapping
 - **Shared File Mapping**
 - POSIX Shared Memory
- ❖ **Synchronization**
 - POSIX Semaphores

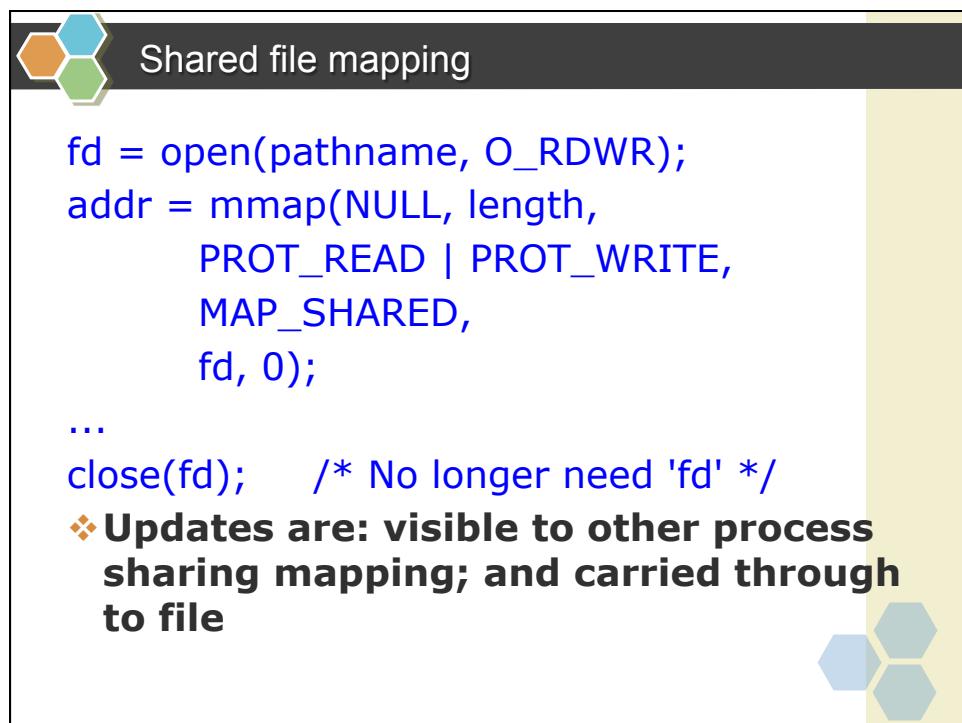
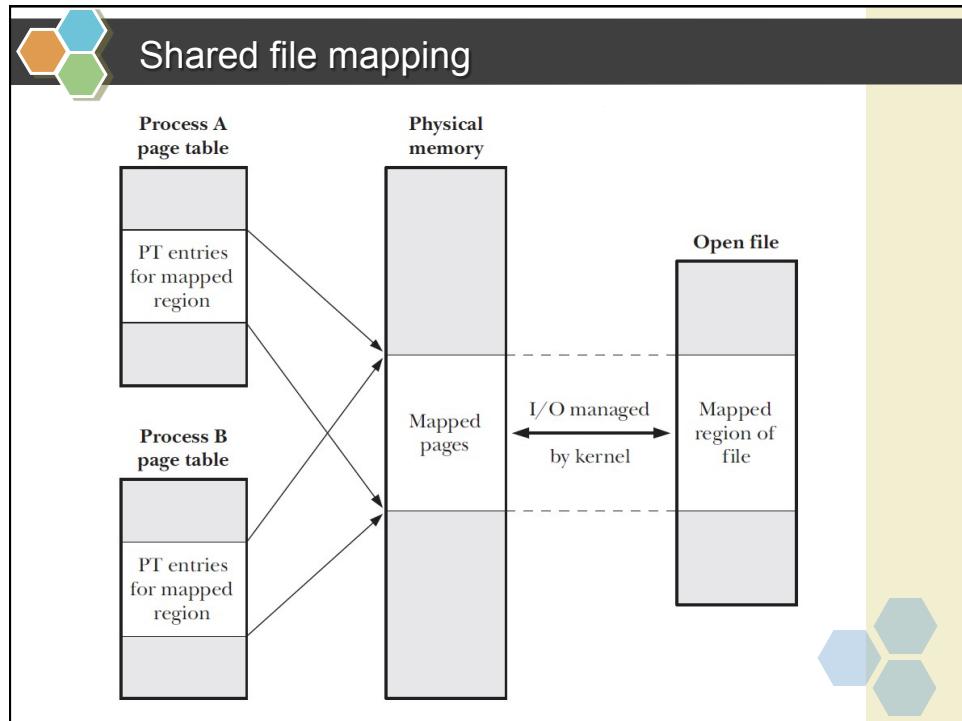
Shared file mapping

- ❖ Share memory between unrelated processes, backed by file
- ❖ `fd = open(...); addr = mmap(..., fd, offset);`

The diagram illustrates the relationship between a process's virtual memory and a file's physical structure. A vertical rectangle labeled 'Process virtual memory' contains a horizontal bar labeled 'mapped region'. An arrow points from the text 'address returned by `mmap()`' to the left edge of this bar. To the right, a horizontal line represents a file with three segments. The first segment is labeled 'offset' with a double-headed arrow indicating its position. The middle segment is labeled 'mapped region' with a double-headed arrow indicating its length.

Shared file mapping

- ❖ `fd = open(...); addr = mmap(..., fd, offset);`
- ❖ Contents of memory initialized from file
- ❖ Updates to memory automatically carried through to file ("memory-mapped I/O")
- ❖ All processes that map same region of file share same memory



Outline

- ❖ **Communication Techniques**
 - Pipes
 - FIFO (named pipes)
 - POSIX Message Queues
 - Shared Memory
 - Shared Anonymous Mapping
 - Shared File Mapping
 - **POSIX Shared Memory**
- ❖ **Synchronization**
 - POSIX Semaphores

POSIX shared memory

- ❖ **Share memory between unrelated process, without creating file in (traditional) filesystem**
 - Don't need to create a file
 - Avoid file I/O overhead

 **POSIX SHM API**

- ❖ Object management
 - `shm_open()`: open/create SHM object
 - `mmap()`: map SHM object
 - `shm_unlink()`: remove SHM object pathname
- ❖ Operations on SHM object via fd returned by `shm_open()`:
 - `fstat()`: retrieve info (size, ownership, permissions)
 - `ftruncate()`: change size
 - `fchown()`: `fchmod()`: change ownership, permissions



 **Opening a POSIX SHM object**

- ❖ `fd = shm_open(name, flags, mode);`
- ❖ Open+create new / open existing SHM object
- ❖ name has form `/somename`
 - Can be seen in dedicated tmpfs at `/dev/shm`
- ❖ Returns fd, a file descriptor
 - Used by rest of API



 Opening a POSIX SHM object

- ❖ `fd = shm_open(name, flags, mode);`
- ❖ flags (analogous to `open()`):
 - `O_CREAT` – create SHM if it doesn't exist
 - `O_EXCL` – create SHM exclusively
 - `O_RDONLY, O_RDWR` – indicates type of access
 - `O_TRUNC` – truncate existing SHM object to zero length
- ❖ mode sets permissions
 - MBZ if `O_CREAT` not specified

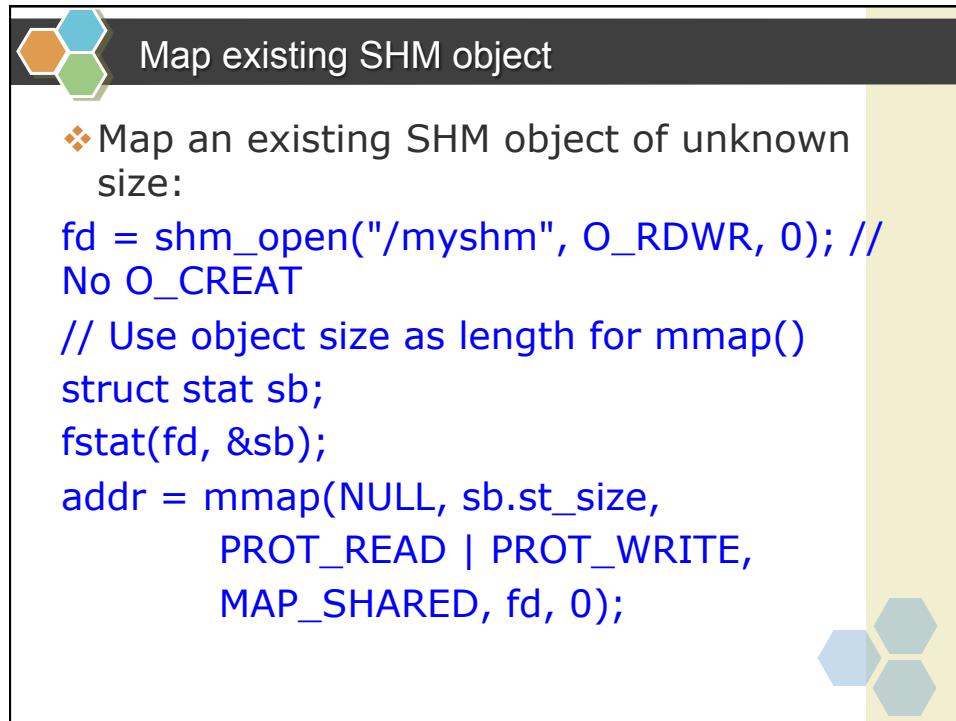


 Create and map new SHM object

- ❖ Create and map a new SHM object of size bytes:

```
fd = shm_open("/myshm",
              O_CREAT | O_EXCL | O_RDWR, 0600);
ftruncate(fd, size); // Set size of object
addr = mmap(NULL, size,
            PROT_READ | PROT_WRITE,
            MAP_SHARED, fd, 0);
```

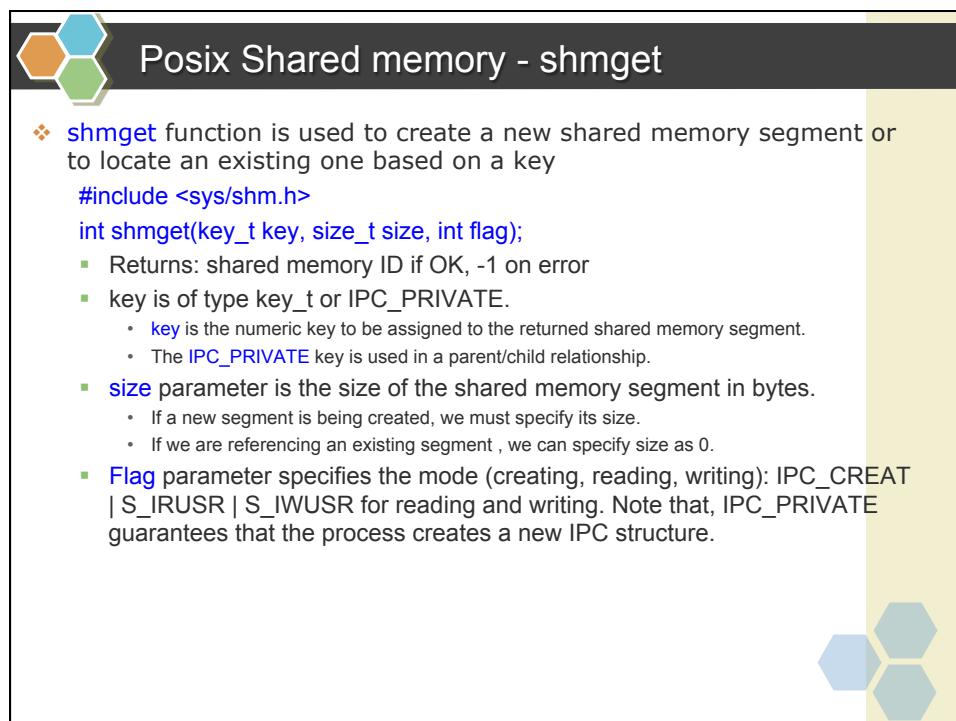




Map existing SHM object

- ❖ Map an existing SHM object of unknown size:

```
fd = shm_open("/myshm", O_RDWR, 0); //  
No O_CREAT  
  
// Use object size as length for mmap()  
struct stat sb;  
fstat(fd, &sb);  
addr = mmap(NULL, sb.st_size,  
           PROT_READ | PROT_WRITE,  
           MAP_SHARED, fd, 0);
```



Posix Shared memory - shmget

- ❖ **shmget** function is used to create a new shared memory segment or to locate an existing one based on a key

```
#include <sys/shm.h>  
int shmget(key_t key, size_t size, int flag);
```

- Returns: shared memory ID if OK, -1 on error
- key is of type key_t or IPC_PRIVATE.
 - **key** is the numeric key to be assigned to the returned shared memory segment.
 - The **IPC_PRIVATE** key is used in a parent/child relationship.
- **size** parameter is the size of the shared memory segment in bytes.
 - If a new segment is being created, we must specify its size.
 - If we are referencing an existing segment , we can specify size as 0.
- **Flag** parameter specifies the mode (creating, reading, writing): IPC_CREAT | S_IRUSR | S_IWUSR for reading and writing. Note that, IPC_PRIVATE guarantees that the process creates a new IPC structure.

Posix Shared memory - shmat

- ❖ Once a shared memory segment has been created, a process attaches it to its address space by calling shmat.

```
#include <sys/shm.h>
void *shmat(int shmid, const void *addr, int flag);
```

- **shmid**: shared memory ID i.e., the return value of shmget
- **addr**: address in the calling process at which the segment is attached depends on the addr argument. Set to NULL, so the segment is attached at the first available address selected by the kernel.
- **flag** : If the SHM_RDONLY bit is specified in flag, the segment is attached read-only. Otherwise (0, or NULL), the segment is attached read/write.
- Returned value: the address at which the segment is attached, or -1 on error occur.

- ❖ kernel maintains a structure for each shared segment (e.g., number of current attaches, pid of creator, size of segment in bytes, last attach time, etc.)

Posix Shared Memory

- ❖ When we're done with a shared memory segment, we call shmdt to detach it.
- ❖ shmdt does not remove the identifier and its associated data structure from the system.

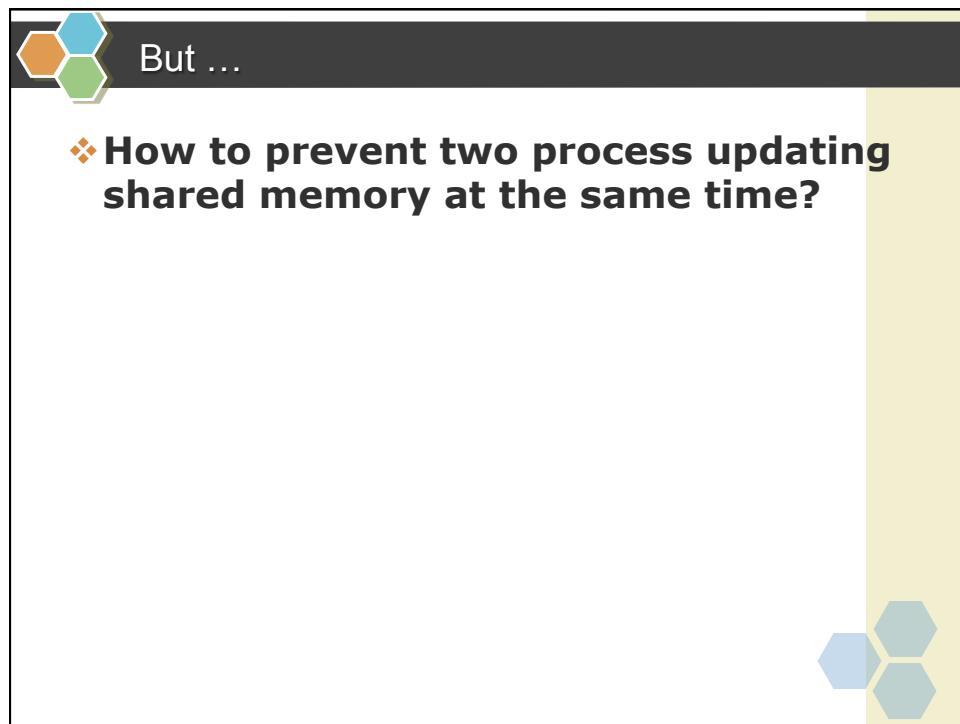
```
#include <sys/shm.h>
int shmdt(void *addr);
```

- **addr**: is the value that was returned by a previous call to shmat.
Returned value: 0 if OK, -1 on error

- ❖ The identifier remains in existence until some process specifically removes it by calling shmctl with a command of IPC_RMID.

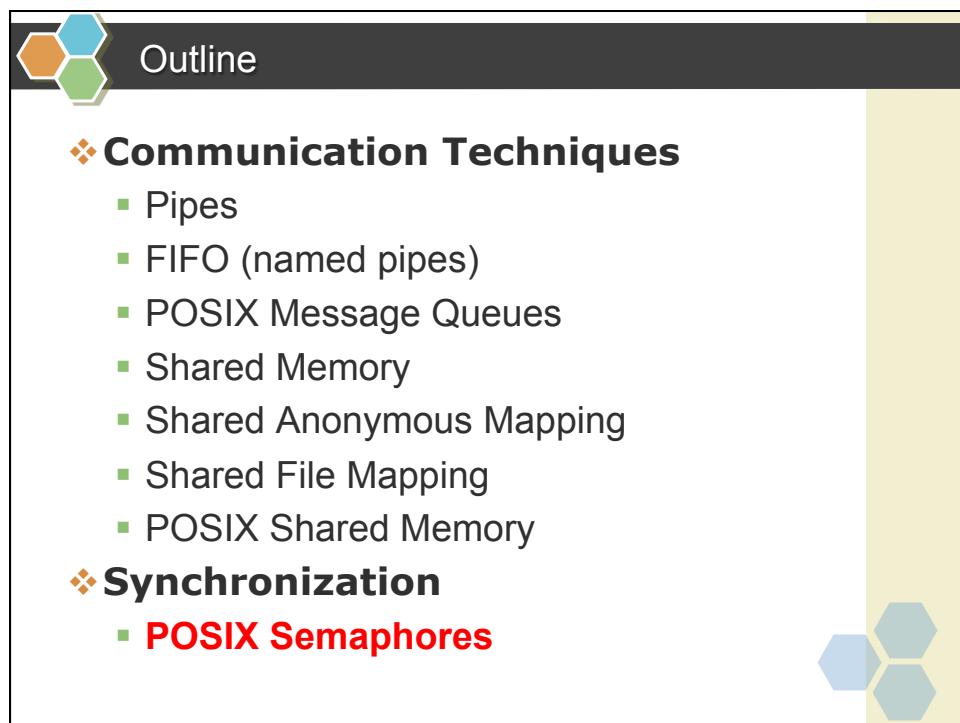
```
#include <sys/shm.h>
int shmctl(int shmid, IPC_RMID, NULL);
```

- ❖



But ...

❖ **How to prevent two process updating shared memory at the same time?**



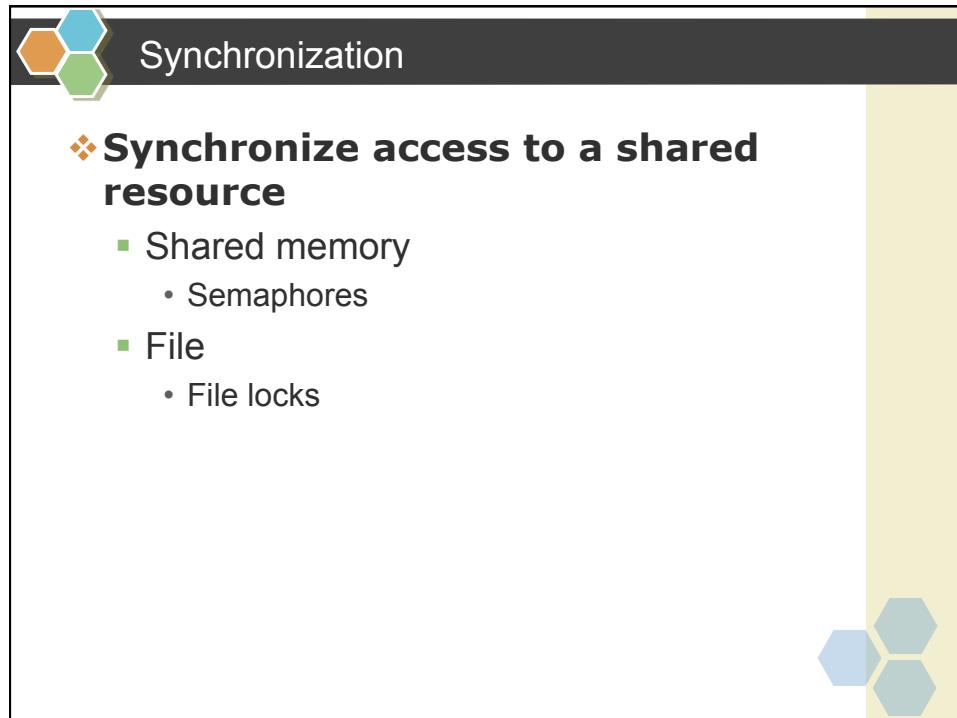
Outline

❖ **Communication Techniques**

- Pipes
- FIFO (named pipes)
- POSIX Message Queues
- Shared Memory
- Shared Anonymous Mapping
- Shared File Mapping
- POSIX Shared Memory

❖ **Synchronization**

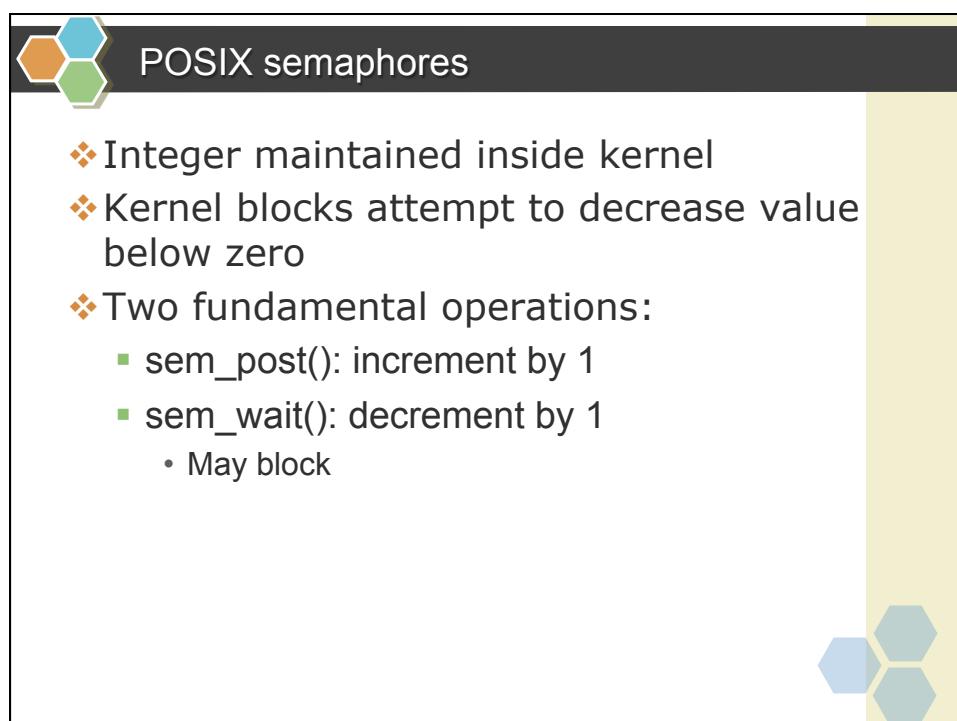
- **POSIX Semaphores**



Synchronization

❖ **Synchronize access to a shared resource**

- Shared memory
 - Semaphores
- File
 - File locks



POSIX semaphores

- ❖ Integer maintained inside kernel
- ❖ Kernel blocks attempt to decrease value below zero
- ❖ Two fundamental operations:
 - `sem_post()`: increment by 1
 - `sem_wait()`: decrement by 1
 - May block



POSIX semaphores



- ❖ **Semaphore represents a shared resource**
- ❖ **E.g., N shared identical resources ==> initial value of semaphore is N**
- ❖ **Common use: binary value**
 - Single resource (e.g., shared memory)



Unnamed and named semaphores



- ❖ **Two types of POSIX semaphore:**
 - Unnamed
 - Embedded in shared memory
 - Named
 - Independent, named objects

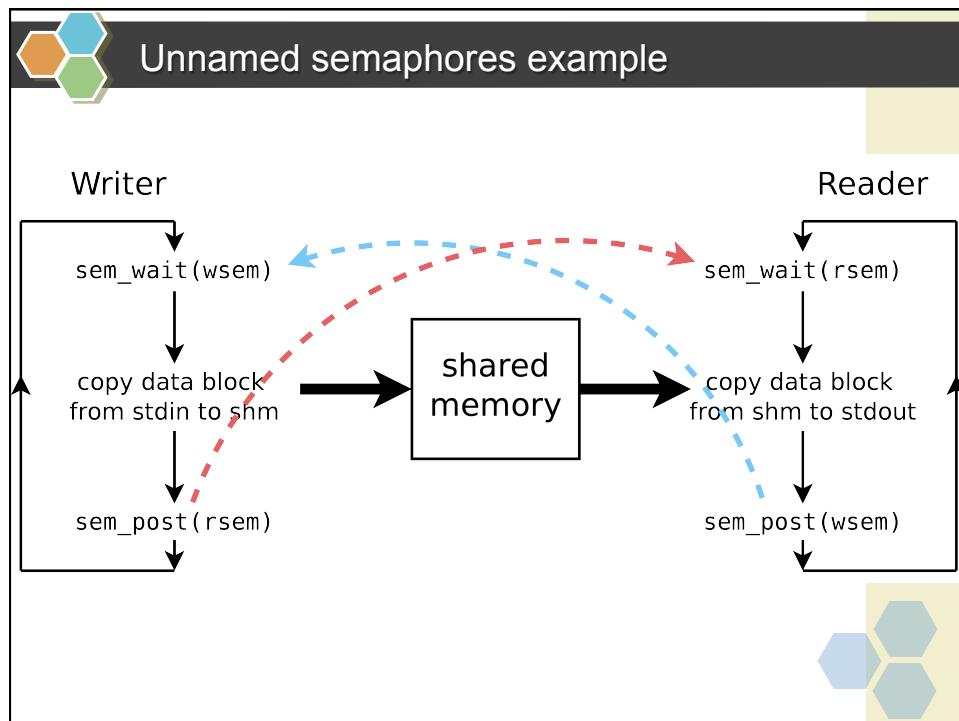


Unnamed semaphores API

- ❖ **sem_init(semp, pshared, value)**: initialize semaphore pointed to by semp to value
 - `sem_t *semp`
 - `pshared`: 0, thread sharing; != 0, process sharing
- ❖ **sem_post(semp)**: add 1 to value
sem_wait(semp): subtract 1 from value
- ❖ **sem_destroy(semp)**: free semaphore, release resources back to system
 - Must be no waiters!

Unnamed semaphores example

- ❖ Two processes, writer and reader
- ❖ Sending data through POSIX shared memory
- ❖ Two unnamed POSIX semaphores inside shm enforce alternating access to shm



Header file

```

#define BUF_SIZE 1024

struct shmbuf {
    sem_t wsem;           // Writer semaphore
    sem_t rsem;           // Reader semaphore
    int cnt;              // Number of bytes used in 'buf'
    char buf[BUF_SIZE];   // Data being transferred
}
  
```



Writer

```

fd = shm_open(SHM_PATH, O_CREAT|O_EXCL|O_RDWR,
OBJ_PERMS);
ftruncate(fd, sizeof(struct shmbuf));
shmp = mmap(NULL, sizeof(struct shmbuf),
            PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
sem_init(&shmp->rsem, 1, 0);
sem_init(&shmp->wsem, 1, 1);      // Writer gets first turn
for (xfrs = 0, bytes = 0; ; xfrs++, bytes += shmp->cnt) {
    sem_wait(&shmp->wsem);        // Wait for our turn
    shmp->cnt = read(STDIN_FILENO, shmp->buf, BUF_SIZE);
    sem_post(&shmp->rsem);        // Give reader a turn
    if (shmp->cnt == 0)           // EOF on stdin?
        break;
}
sem_wait(&shmp->wsem);          // Wait for reader to finish

```




Reader

```

fd = shm_open(SHM_PATH, O_RDWR, 0);
shmp = mmap(NULL, sizeof(struct shmbuf),
            PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
for (xfrs = 0, bytes = 0; ; xfrs++) {
    sem_wait(&shmp->rsem);    // Wait for our turn */
    if (shmp->cnt == 0)        // Writer encountered EOF */
        break;
    bytes += shmp->cnt;
    write(STDOUT_FILENO, shmp->buf, shmp->cnt) != shmp-
>cnt);
    sem_post(&shmp->wsem);   // Give writer a turn */
}
sem_post(&shmp->wsem); // Let writer know we're finished

```





Named semaphores API

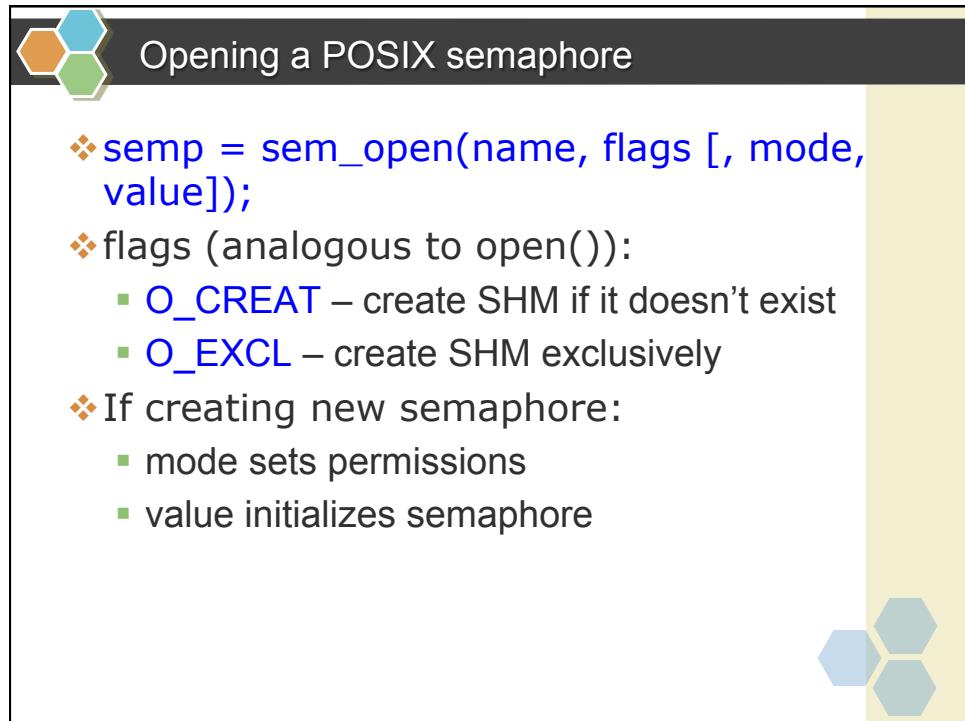
- ❖ **Object management**
 - `sem_open()`: open/create semaphore
 - `sem_unlink()`: remove semaphore pathname



Opening a POSIX semaphore

- ❖ `sem_p = sem_open(name, flags [, mode, value]);`
- ❖ Open+create new / open existing semaphore
- ❖ name has form /somename
 - Can be seen in dedicated tmpfs at /dev/shm
- ❖ Returns `sem_t *`, reference to semaphore
 - Used by rest of API





Opening a POSIX semaphore

- ❖ `sem = sem_open(name, flags [, mode, value]);`
- ❖ flags (analogous to `open()`):
 - `O_CREAT` – create SHM if it doesn't exist
 - `O_EXCL` – create SHM exclusively
- ❖ If creating new semaphore:
 - mode sets permissions
 - value initializes semaphore